

Virtual Memory

'22H2

송 인 식

Outline

- Virtual Memory: Concepts
- Virtual Memory: Implementation

Virtualizing Memory

- Example

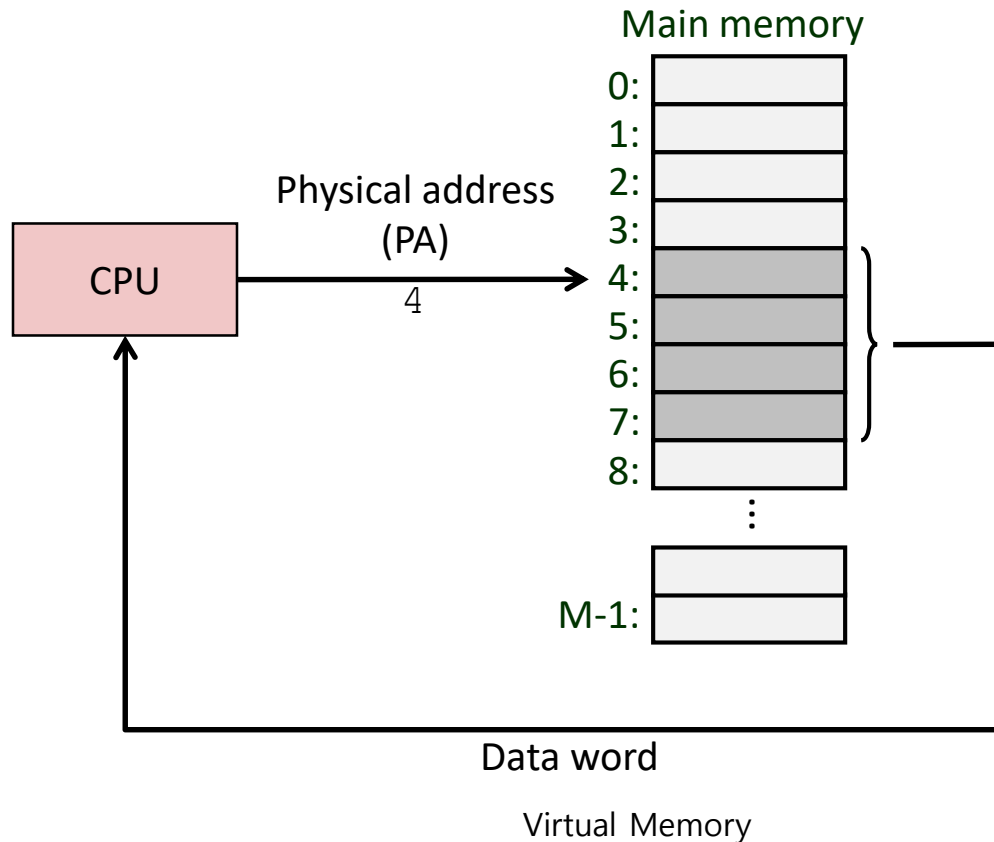
```
#include <stdio.h>
int n = 0;
int main ()
{
    n++;
    printf (“&n = %p, n = %d\n”, &n, n);
}

% ./a.out
&n = 0x0804a024, n = 1
% ./a.out
&n = 0x0804a024, n = 1
```

- What happens if two users simultaneously run this program?

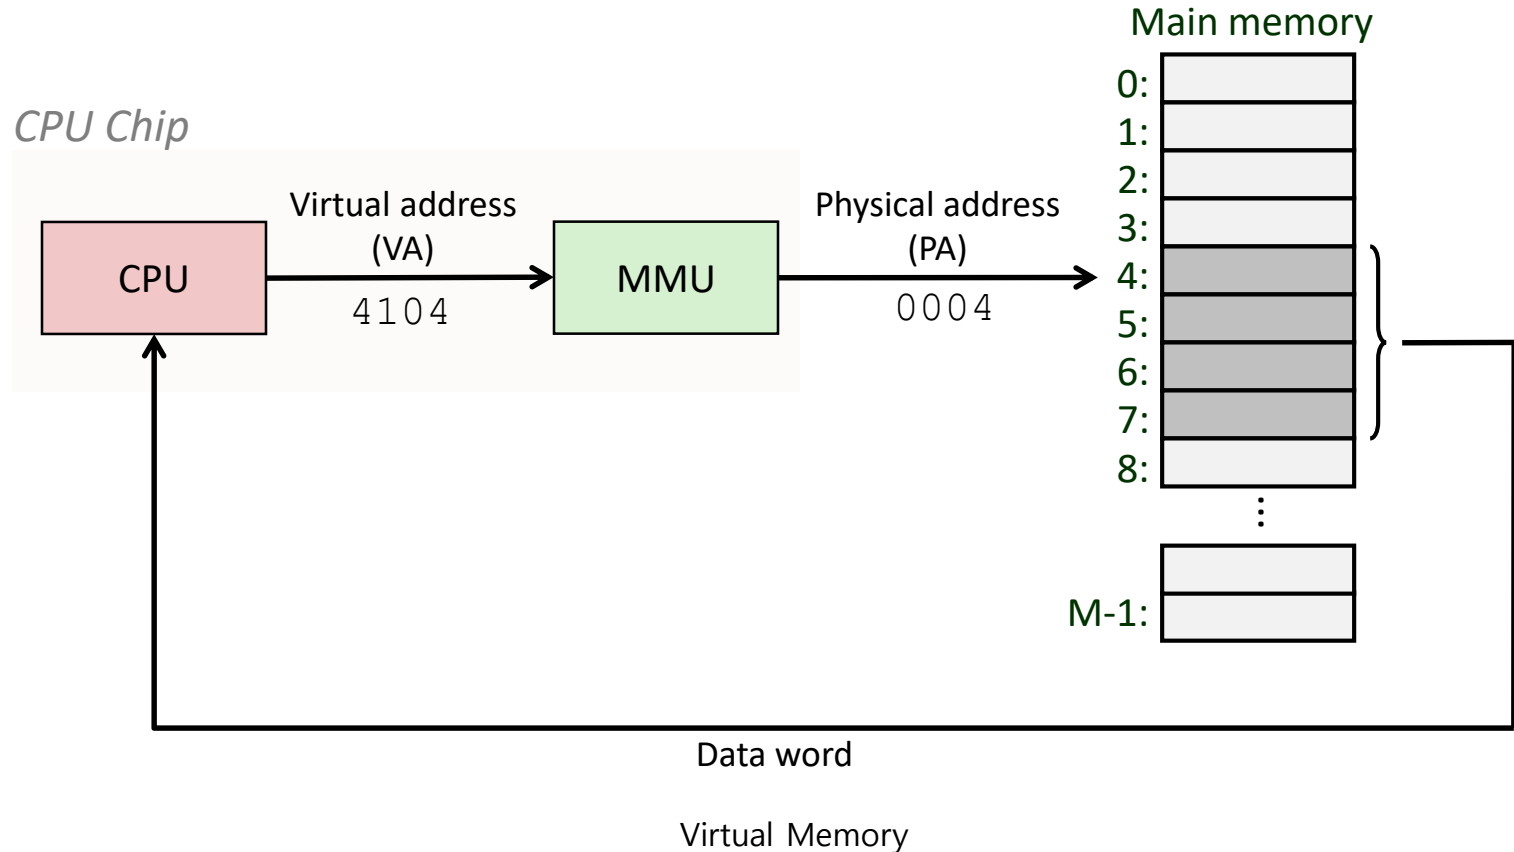
A System Using Physical Addressing

- Used in “simple” systems like embedded microcontrollers in cars, elevators, digital cameras, etc.



A System Using Virtual Addressing

- Used in all modern servers, desktops, and laptops
- One of the great ideas in computer science



Virtual Memory

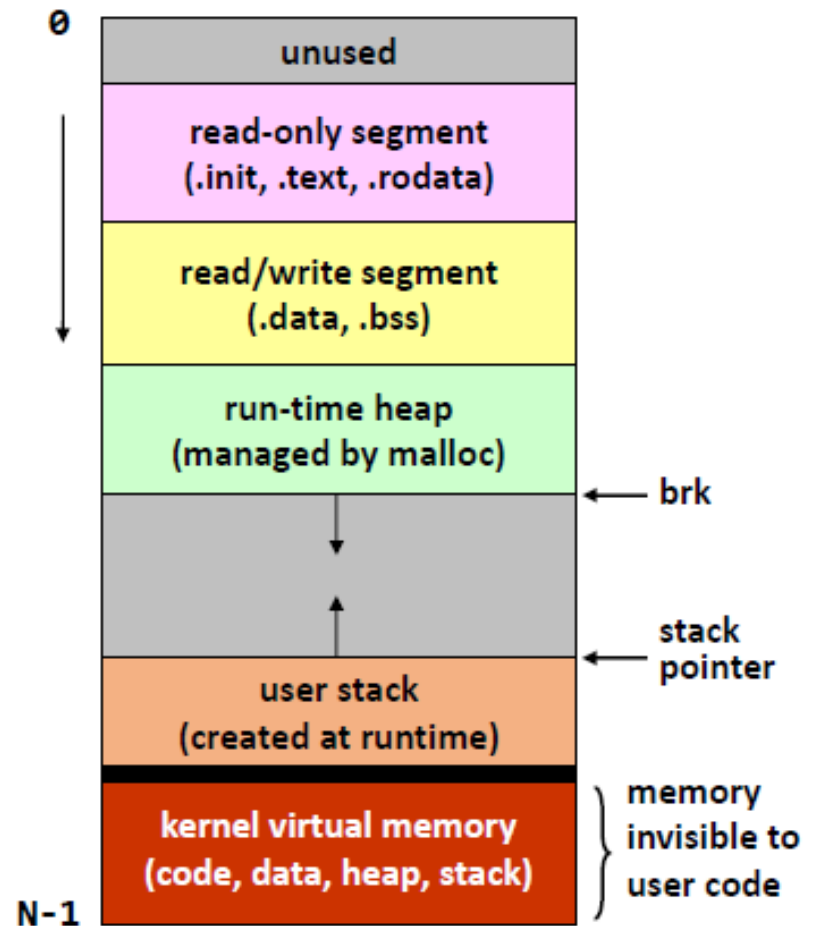
- Each process has its own virtual address space
 - Large and contiguous
 - Use virtual addresses for memory references
 - Virtual addresses are private to each process
- Address translation is performed at run time
 - From a virtual address to the corresponding physical address
- Supports lazy allocation
 - Physical memory is dynamically allocated or released on demand
 - Programs execute without requiring their entire address space to be resident in physical memory

Address Spaces

- **Linear address space:** Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots \}$
- **Virtual address space:** Set of $N = 2^n$ virtual addresses
 $\{0, 1, 2, 3, \dots, N-1\}$
- **Physical address space:** Set of $M = 2^m$ physical addresses
 $\{0, 1, 2, 3, \dots, M-1\}$

(Virtual) Address Space

- Process' abstract view of memory
 - OS provides illusion of private address space to each process
 - Contains all of the memory state of the process
 - Static area
 - Allocated on `exec()`
 - Code & Data
 - Dynamic area
 - Allocated at runtime
 - Can grow or shrink
 - Heap & Stack

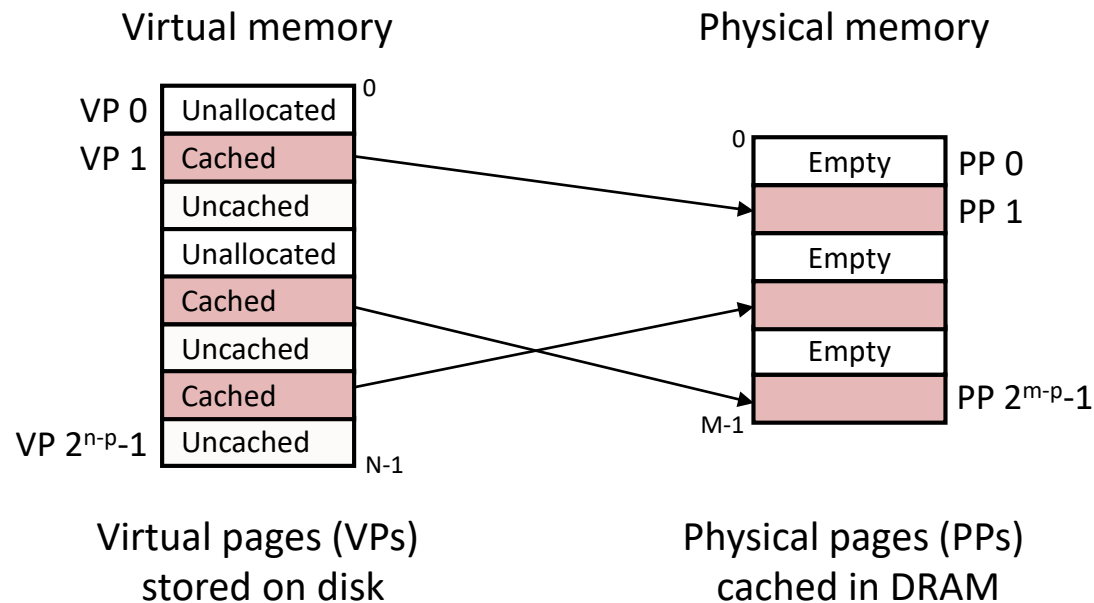


Why Virtual Memory (VM)?

- VM as a tool for caching
 - Uses main memory efficiently
 - Use DRAM as a cache for parts of a virtual address space
- VM as a tool for memory management
 - Simplifies memory management
 - Each process gets the same uniform linear address space
- VM as a tool for memory protection
 - Isolates address spaces
 - One process can't interfere with another's memory
 - User program cannot access privileged kernel information and code

VM as a Tool for Caching

- *Virtual memory* is an array of N contiguous bytes stored on disk.
- The contents of the array on disk are cached in *physical memory* (*DRAM cache*)
 - These cache blocks are called *pages* (size is $P = 2^p$ bytes)

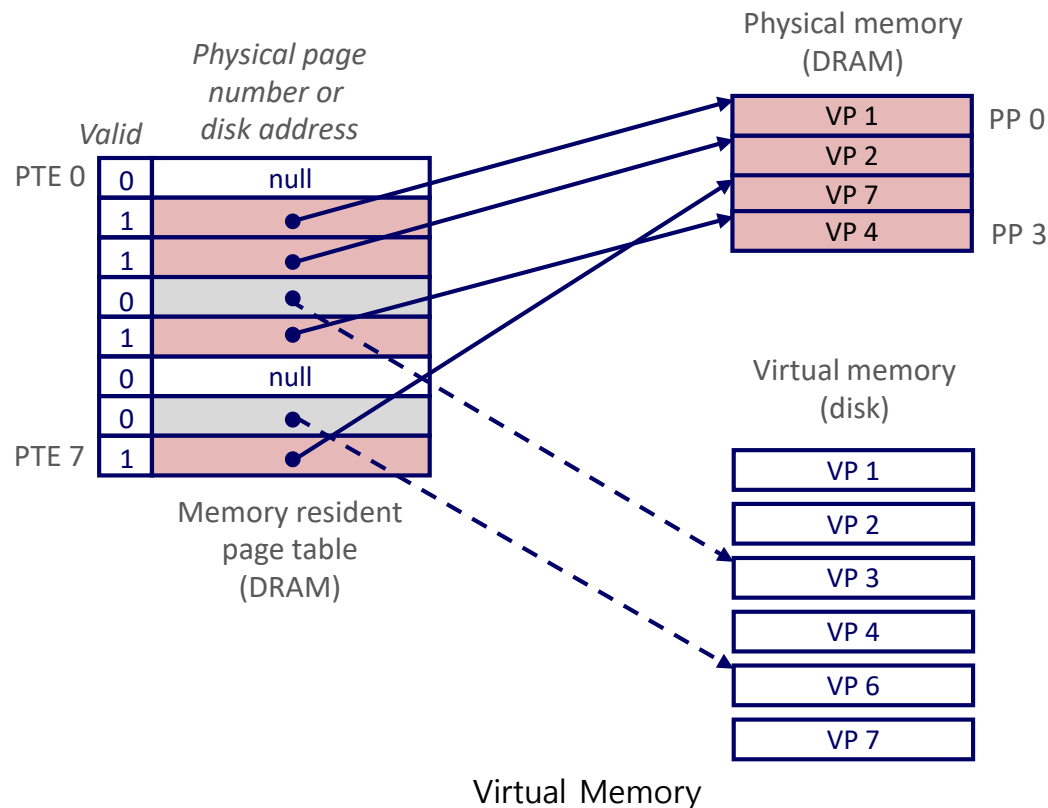


DRAM Cache Organization

- DRAM cache organization driven by the enormous miss penalty
 - DRAM is about **10x** slower than SRAM
 - Disk is about **10,000x** slower than DRAM
- Consequences
 - Large page (block) size: typically 4-8 KB, sometimes 4 MB
 - Fully associative
 - Any VP can be placed in any PP
 - Requires a “large” mapping function – different from CPU caches
 - Highly sophisticated, expensive replacement algorithms
 - Too complicated and open-ended to be implemented in hardware
 - Write-back rather than write-through

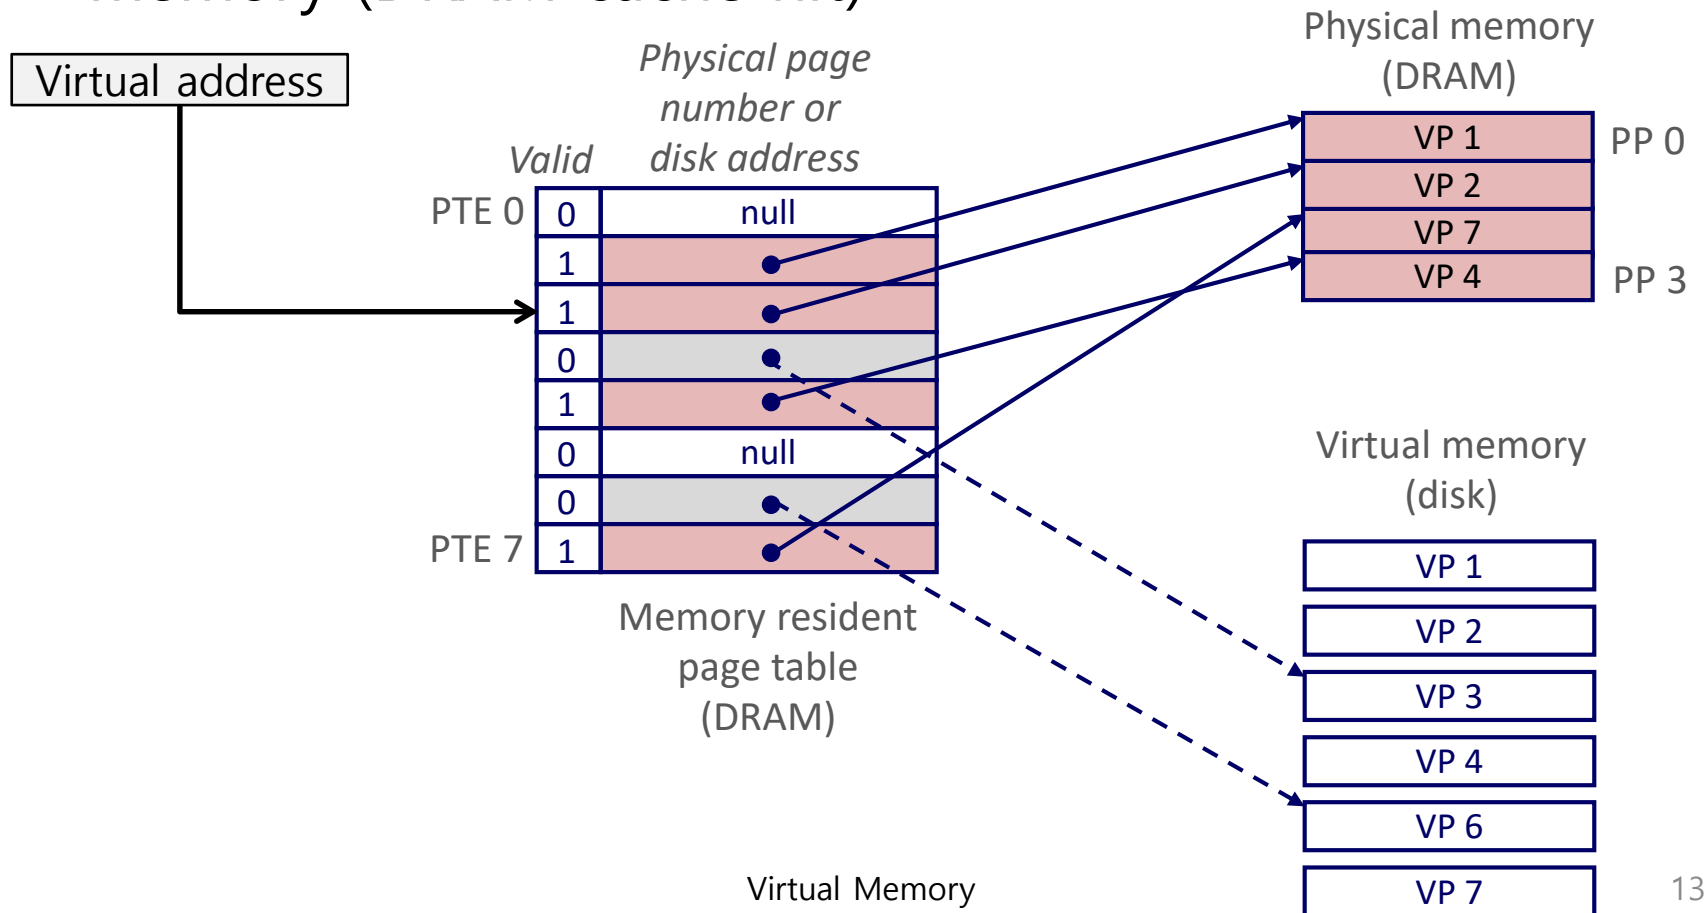
Enabling data structure: Page Table

- A *page table* is an array of page table entries (PTEs) that maps virtual pages to physical pages.
 - Per-process kernel data structure in DRAM



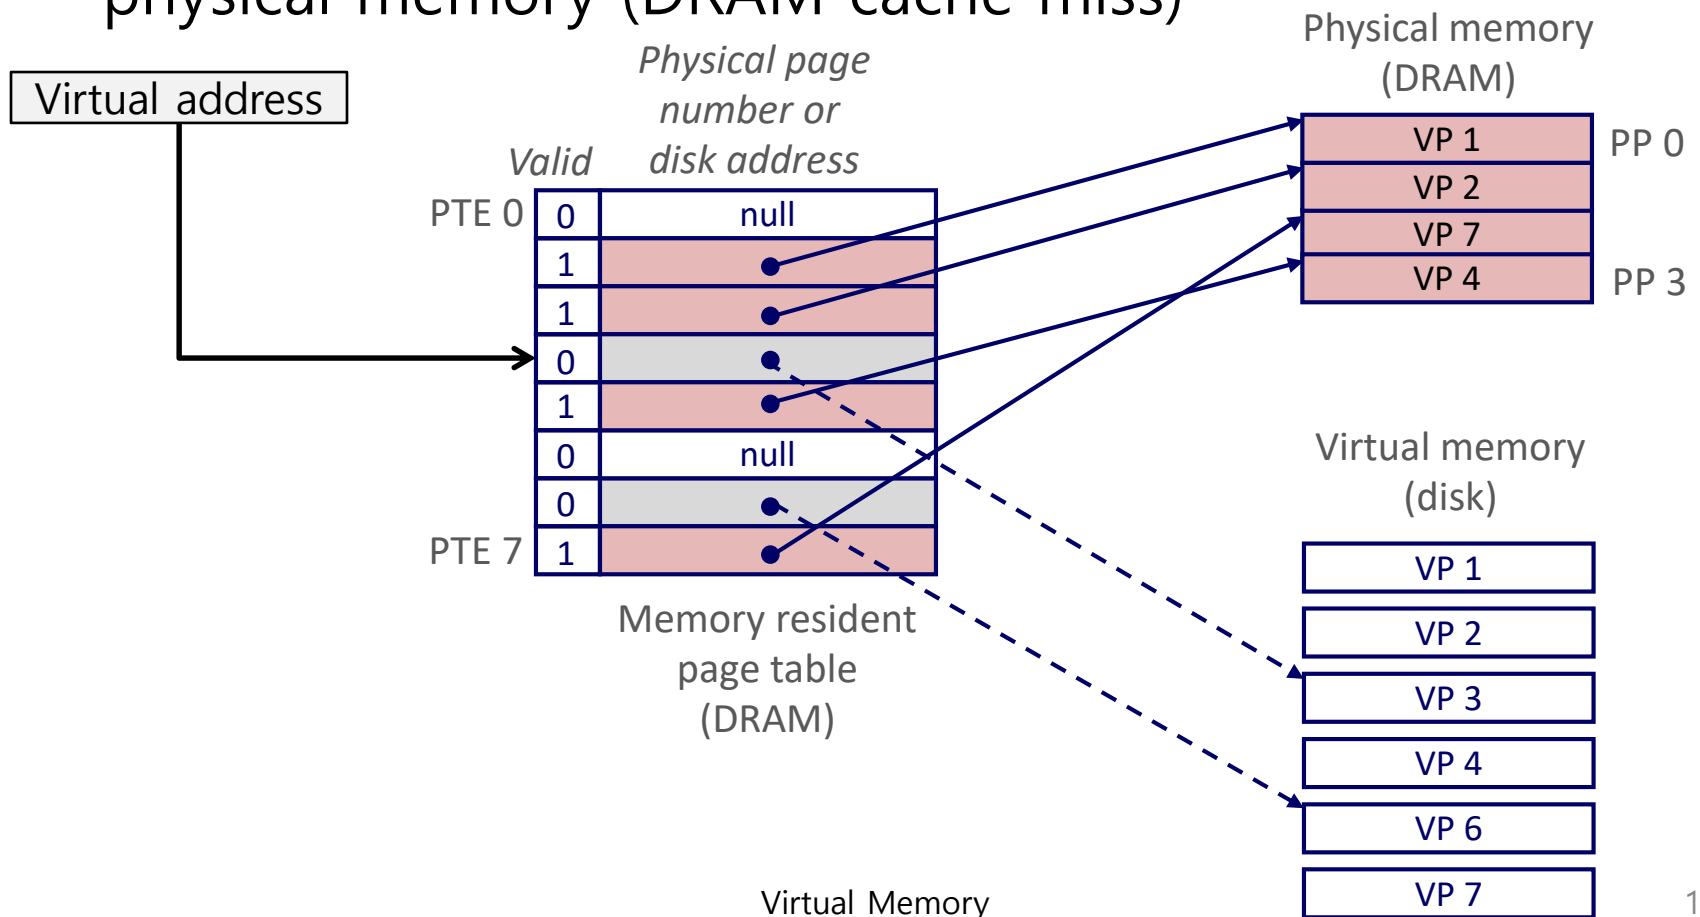
Page Hit

- *Page hit*: reference to VM word that is in physical memory (DRAM cache hit)



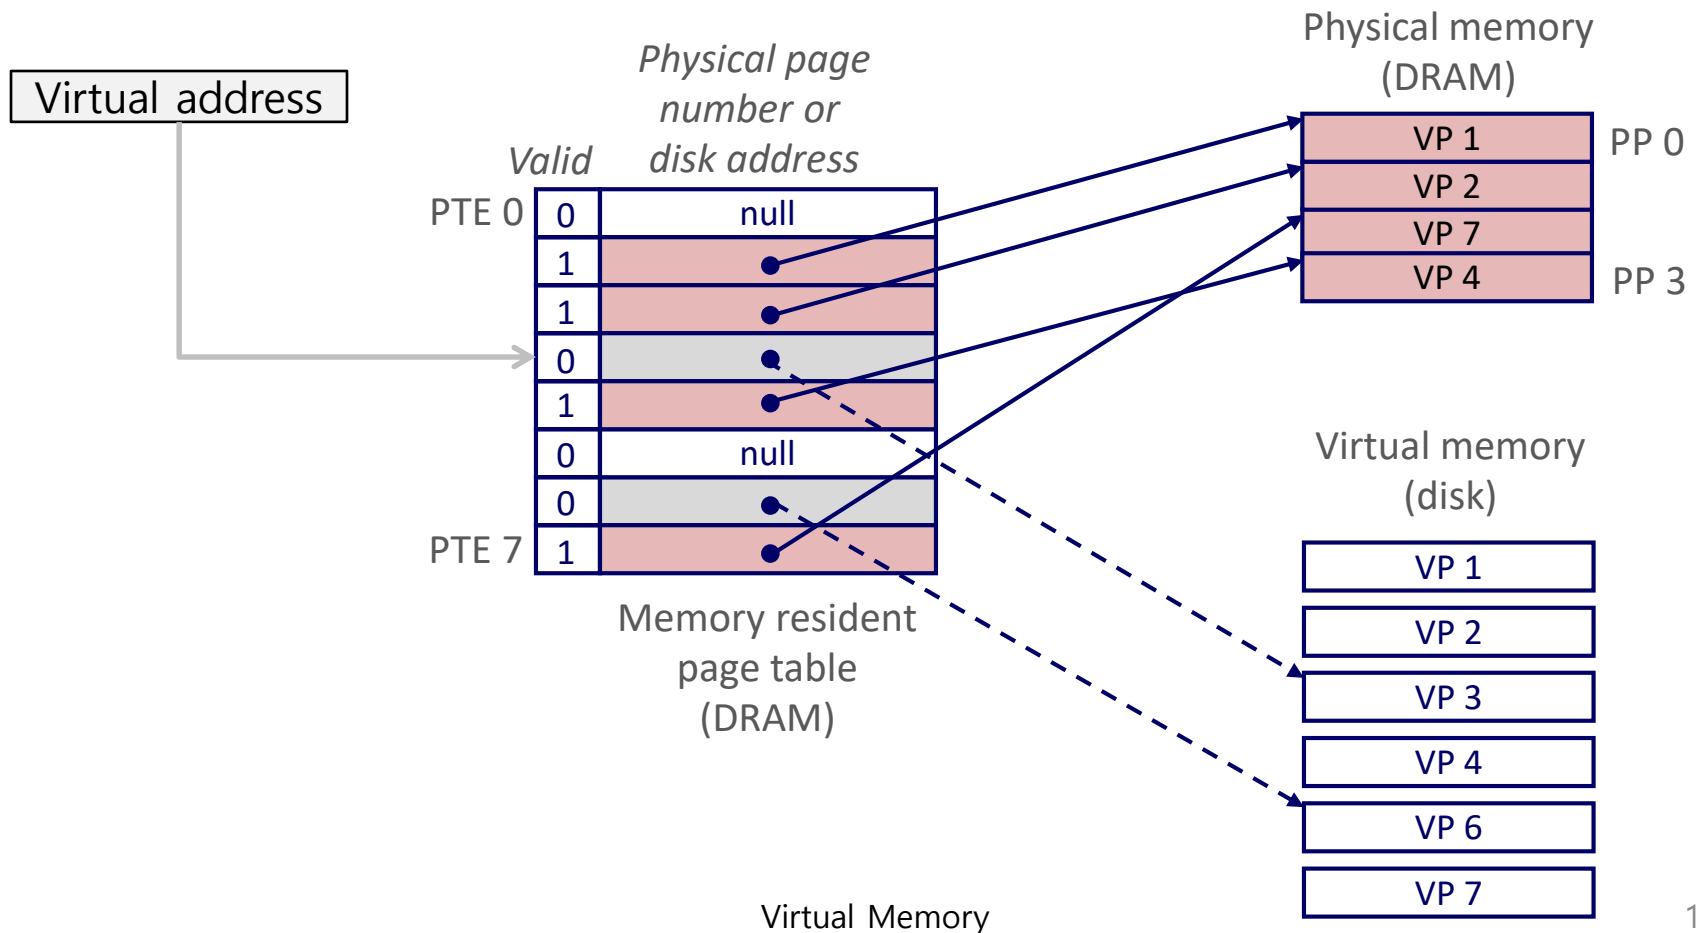
Page Fault

- *Page fault*: reference to VM word that is not in physical memory (DRAM cache miss)



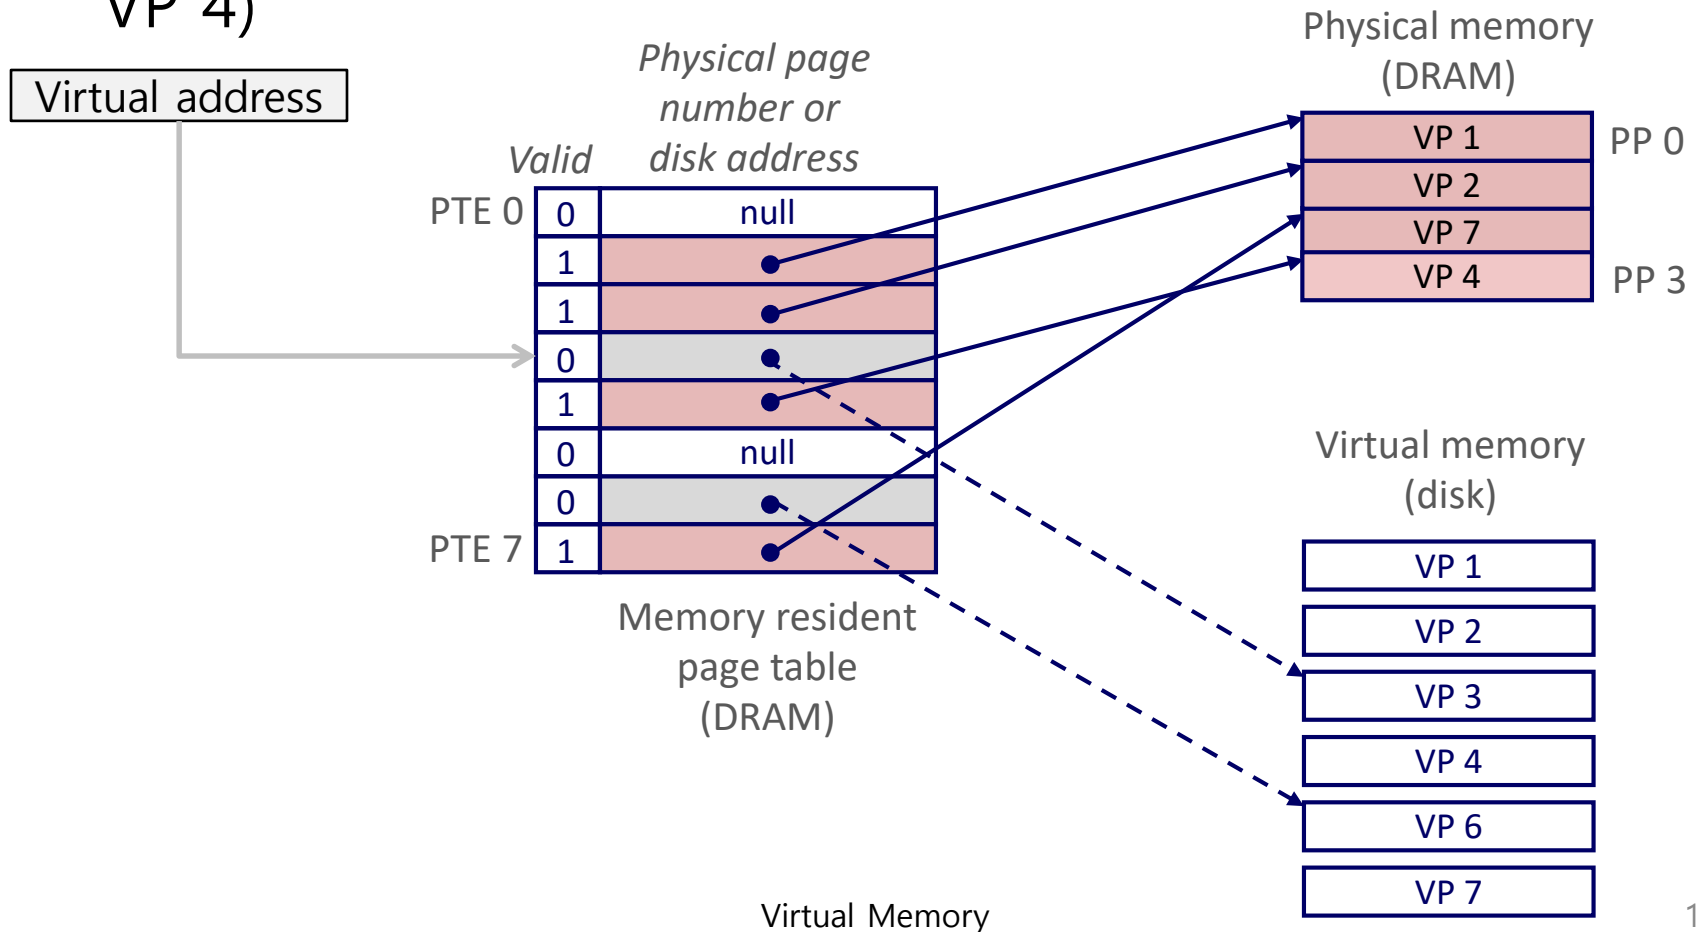
Handling Page Fault #1

- Page miss causes page fault (an exception)



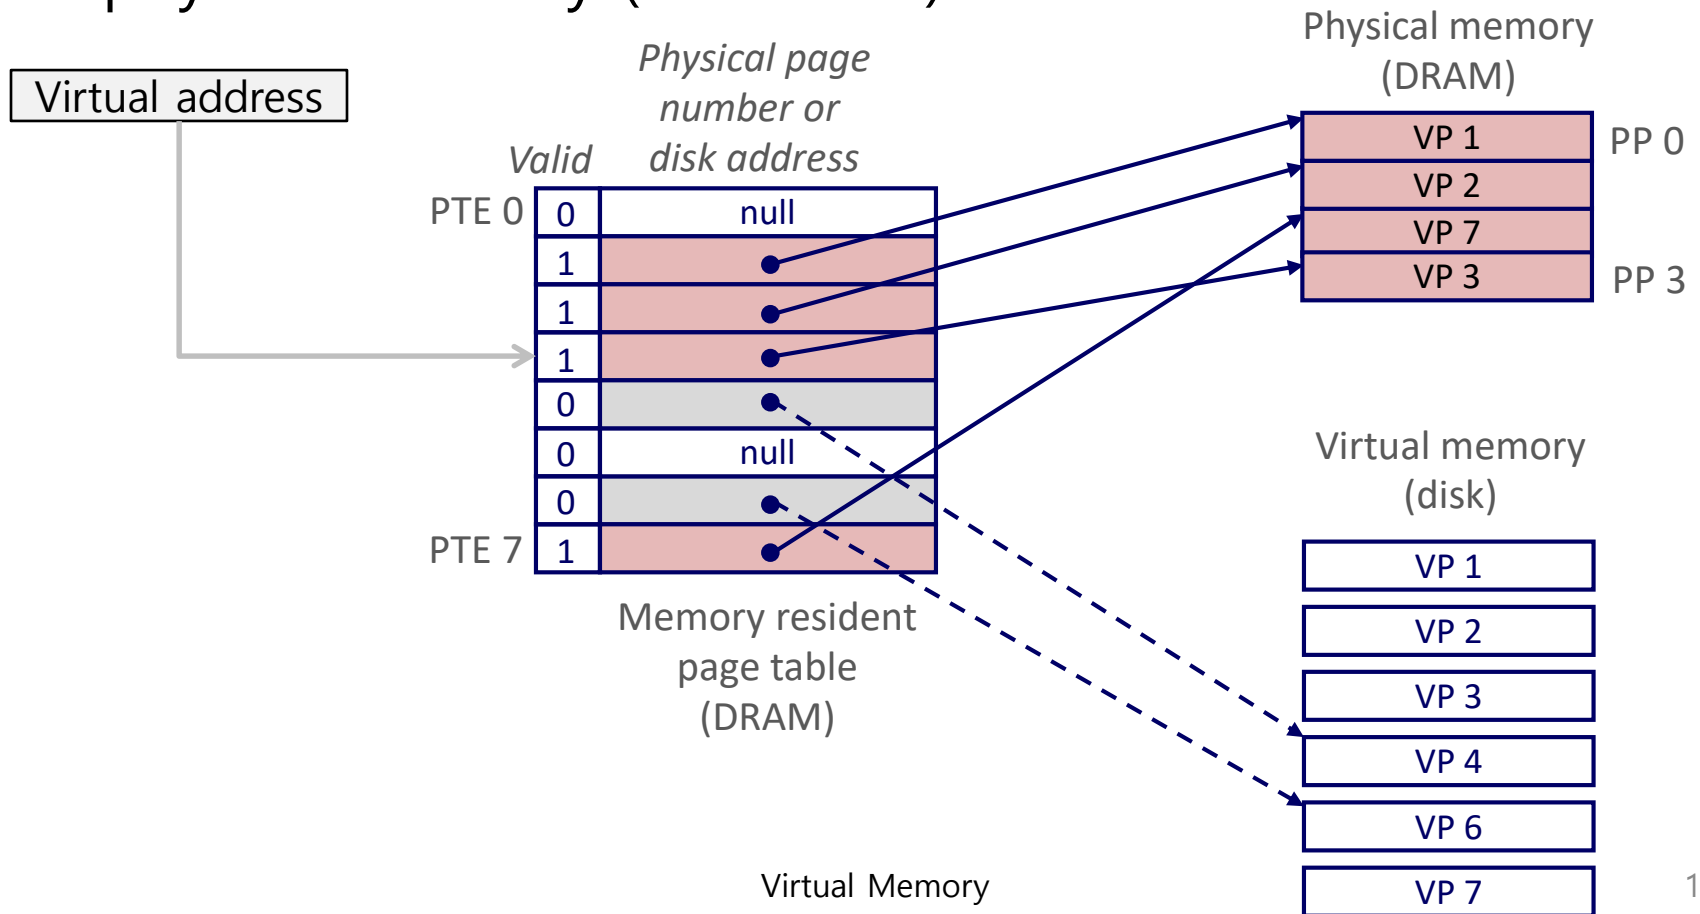
Handling Page Fault #2

- Page fault handler selects a victim to be evicted (here VP 4)



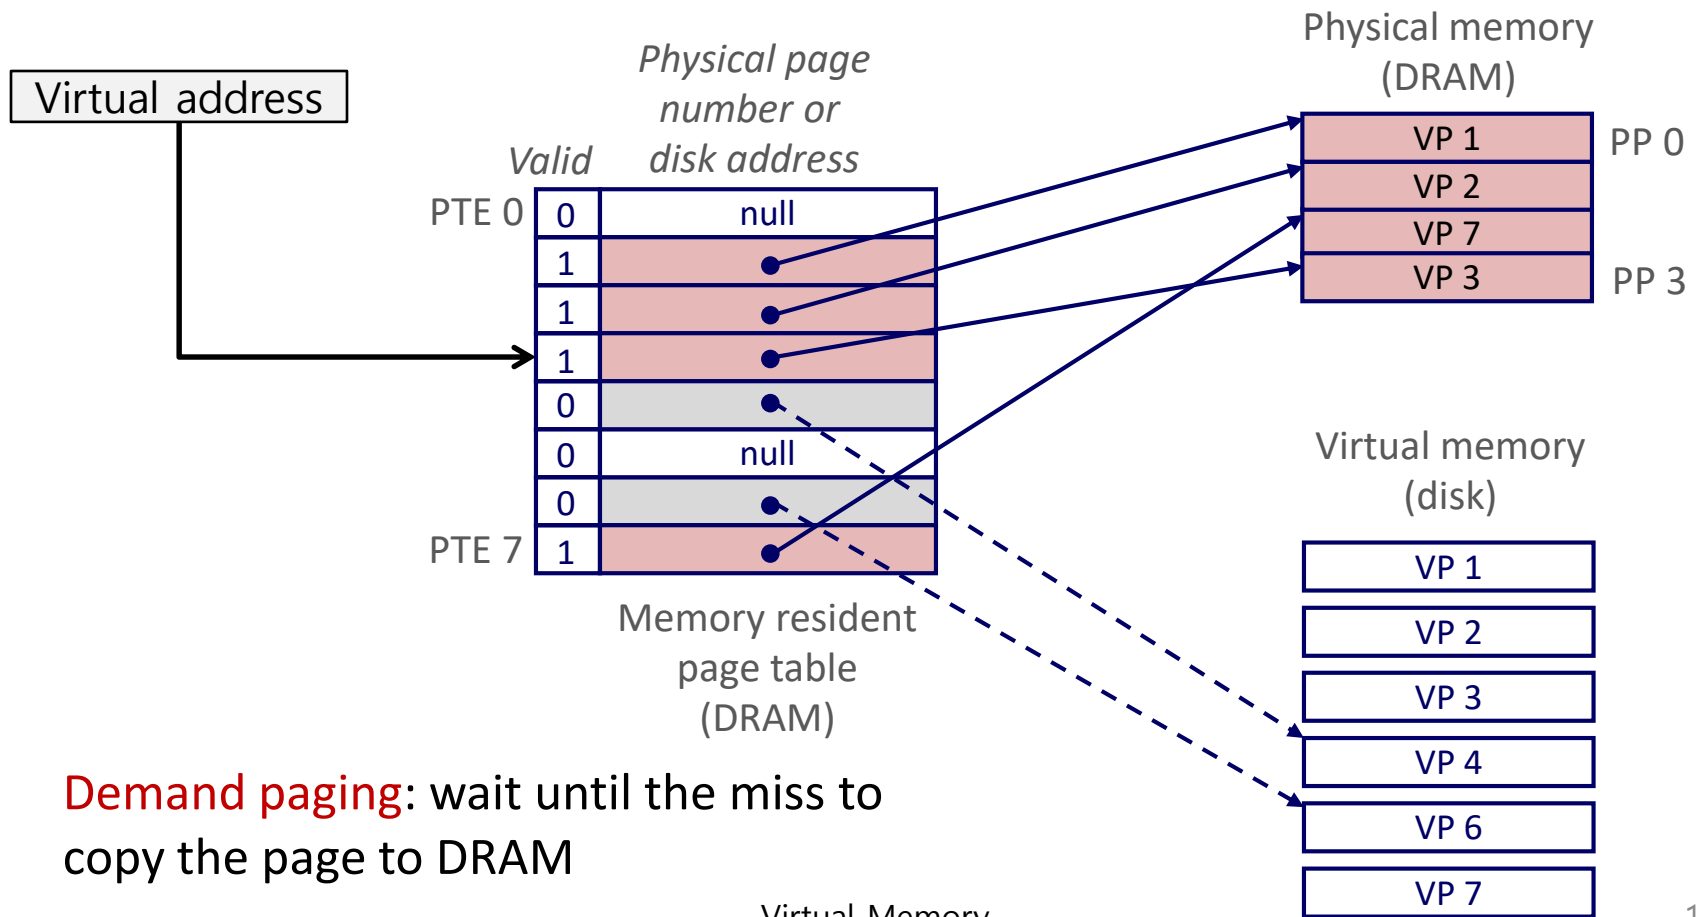
Handling Page Fault #3

- Page fault handler loads the faulting page into physical memory (here VP 3)



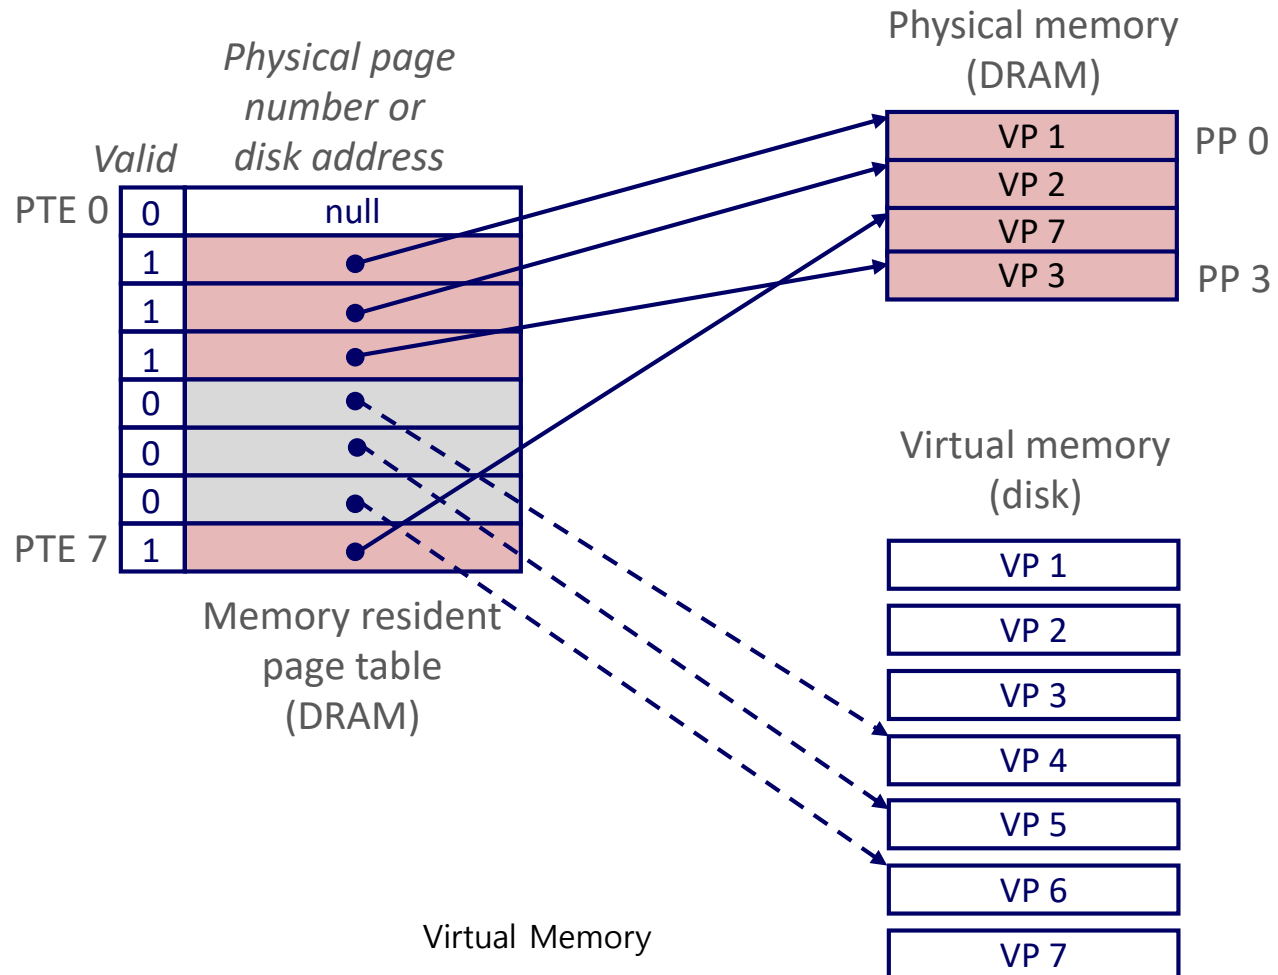
Handling Page Fault #4

- Offending instruction is restarted: page hit!



Allocating Pages

- Allocating a new page (VP 5) of virtual memory.




Locality to the Rescue Again!

- Virtual memory works because of locality
- At any point in time, programs tend to access a set of active virtual pages called the *working set*
 - Programs with better temporal locality will have smaller working sets
- If (working set size < main memory size)
 - Good performance for one process after compulsory misses
- If (SUM(working set sizes) > main memory size)
 - *Thrashing*: Performance meltdown where pages are moved (copied) in and out continuously


VM as a Tool for Memory Management

- Key idea: each process has its own virtual address space
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory
 - Well chosen mappings simplify memory allocation and management



VM as a Tool for Memory Management

- Memory allocation
 - Each virtual page can be mapped to any physical page
 - A virtual page can be stored in different physical pages at different times
- Sharing code and data among processes
 - Map multiple virtual pages to the same physical page (here: PP 6)



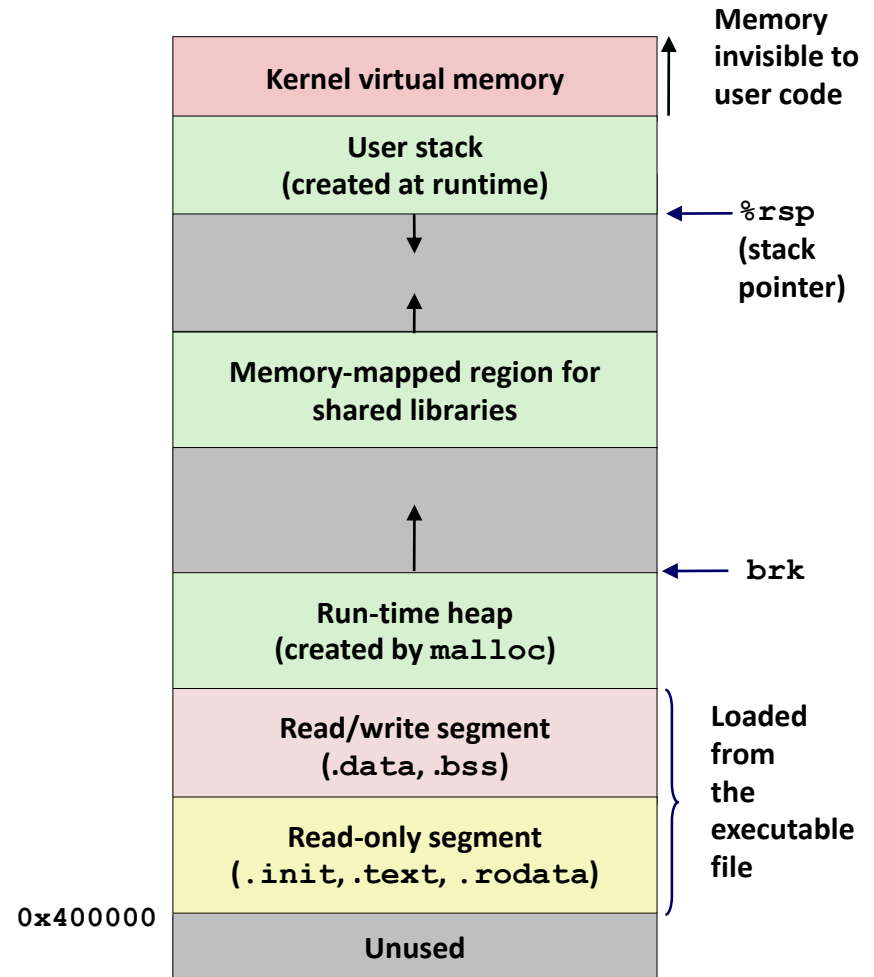
Simplifying Linking and Loading

- Linking

- Each program has similar virtual address space
- Code, stack, and shared libraries always start at the same address

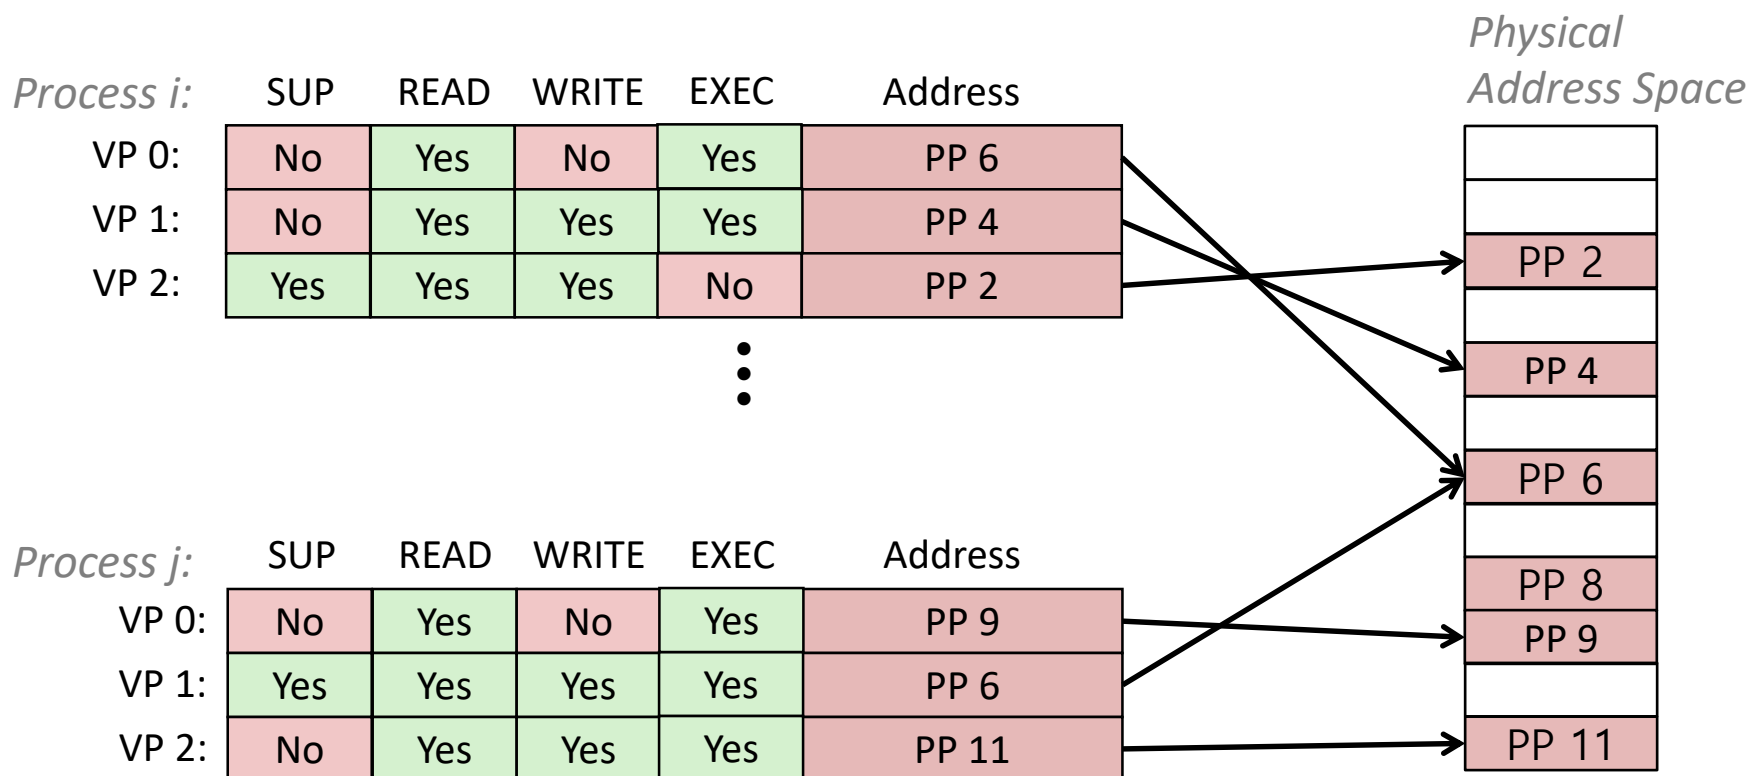
- Loading

- **execve()** allocates virtual pages for **.text** and **.data** sections = creates PTEs marked as invalid
- The **.text** and **.data** sections are copied, page by page, on demand by the virtual memory system



VM as a Tool for Memory Protection

- Extend PTEs with permission bits
- MMU checks these bits on each access



Outline

- Virtual Memory: Concepts
- Virtual Memory: Implementation

VM Address Translation

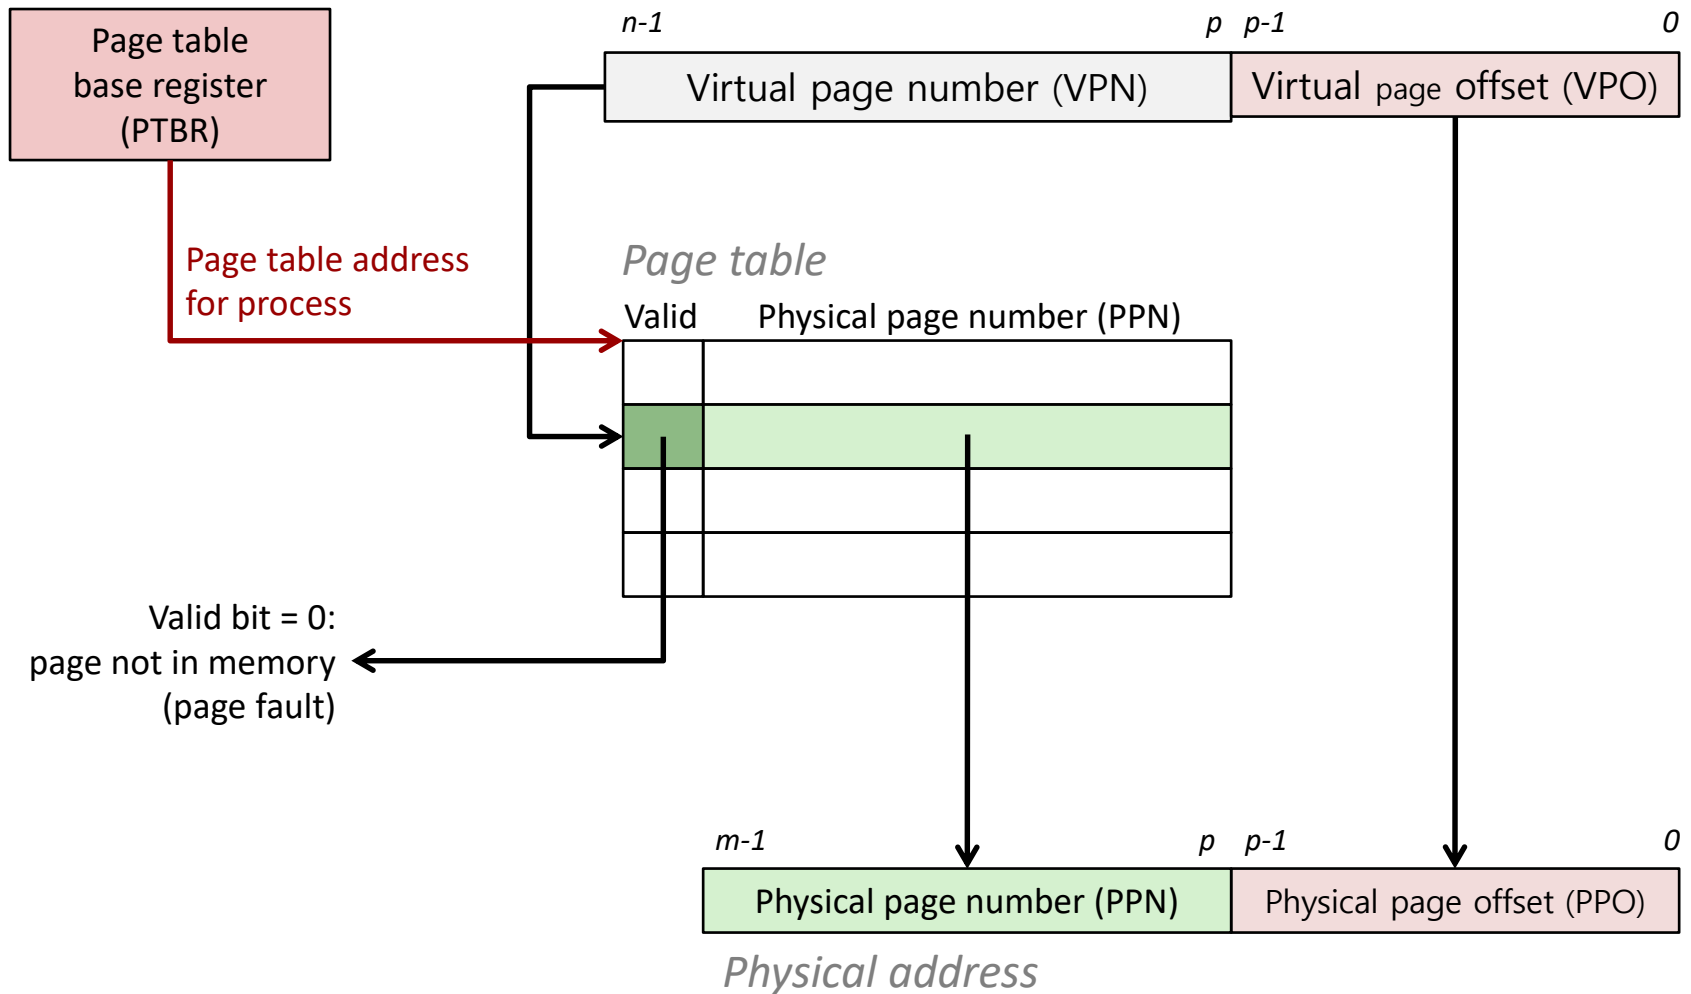
- Virtual Address Space
 - $V = \{0, 1, \dots, N-1\}$
- Physical Address Space
 - $P = \{0, 1, \dots, M-1\}$
- Address Translation
 - ***MAP***: $V \rightarrow P \cup \{\emptyset\}$
 - For virtual address a :
 - ***MAP***(a) = a' if data at virtual address a is at physical address a' in P
 - ***MAP***(a) = \emptyset if data at virtual address a is not in physical memory
 - Either invalid or stored on disk

Summary of Address Translation Symbols

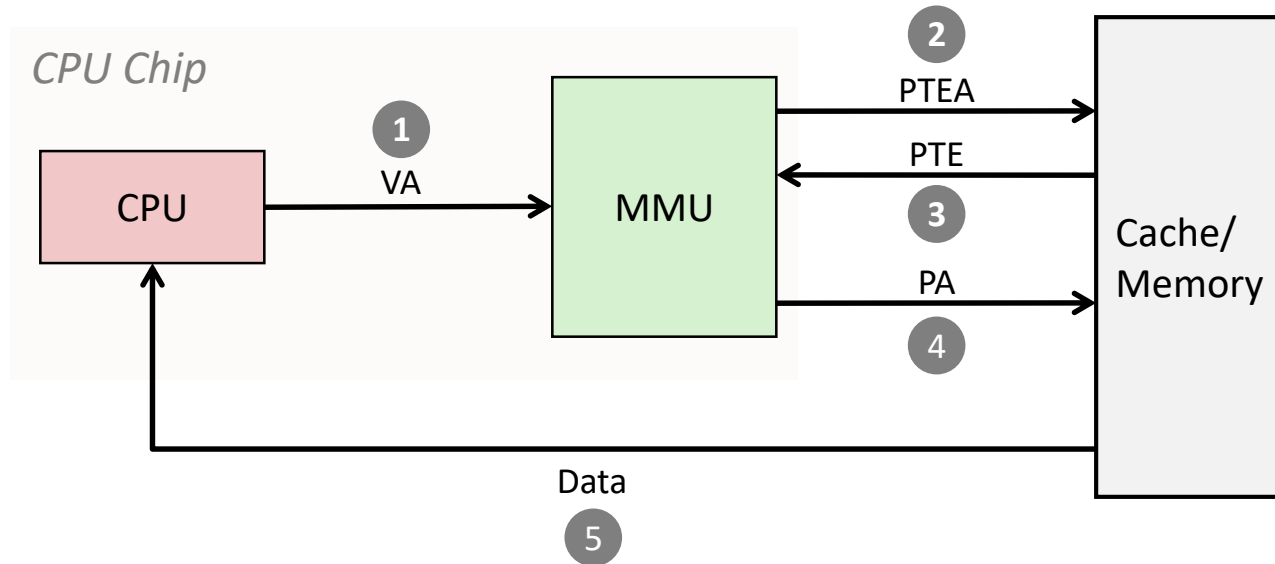
- Basic Parameters
 - **N** = 2^n : Number of addresses in virtual address space
 - **M** = 2^m : Number of addresses in physical address space
 - **P** = 2^p : Page size (bytes)
- Components of the virtual address (VA)
 - **VPO**: Virtual page offset
 - **VPN**: Virtual page number
 - **TLBI**: TLB index
 - **TLBT**: TLB tag
- Components of the physical address (PA)
 - **PPO**: Physical page offset (same as VPO)
 - **PPN**: Physical page number
 - **CO**: Byte offset within cache line
 - **CI**: Cache index
 - **CT**: Cache tag

Address Translation With a Page Table

Virtual address

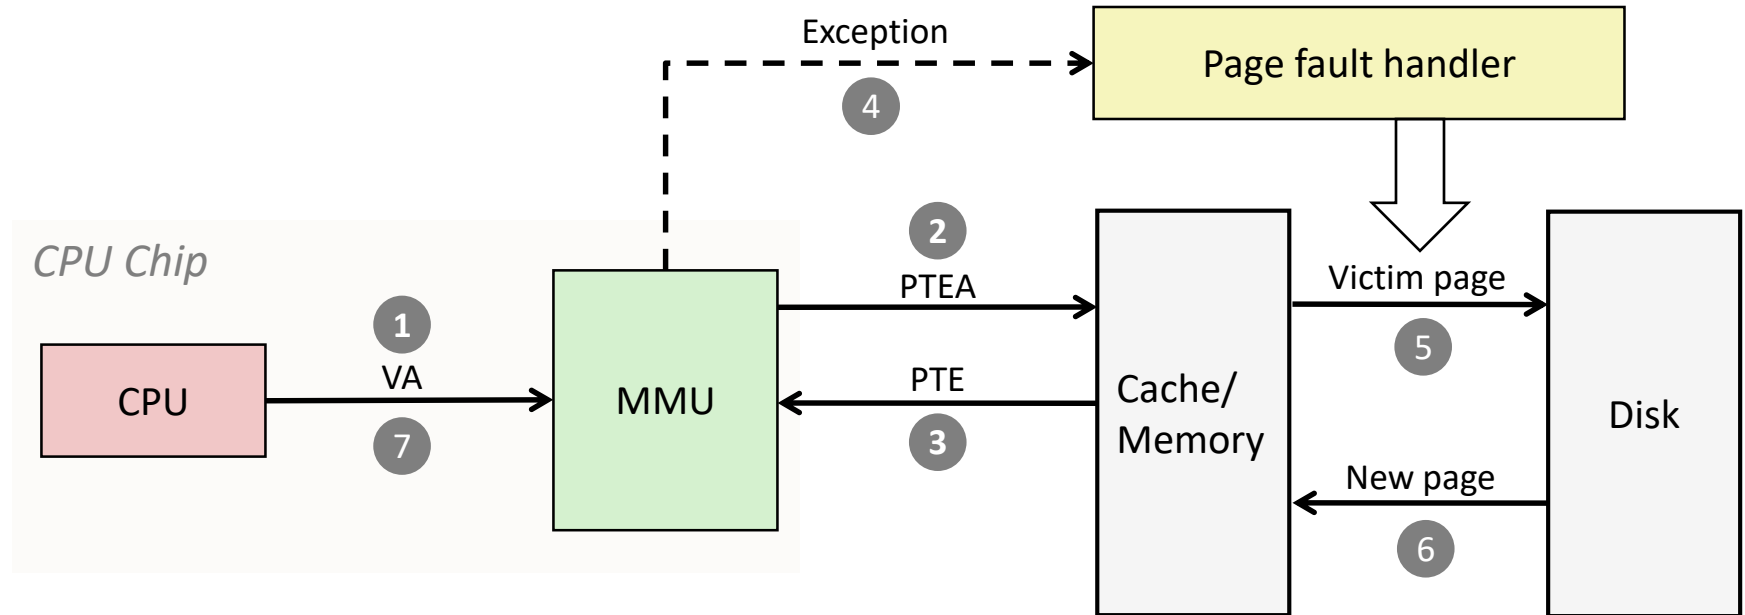


Address Translation: Page Hit



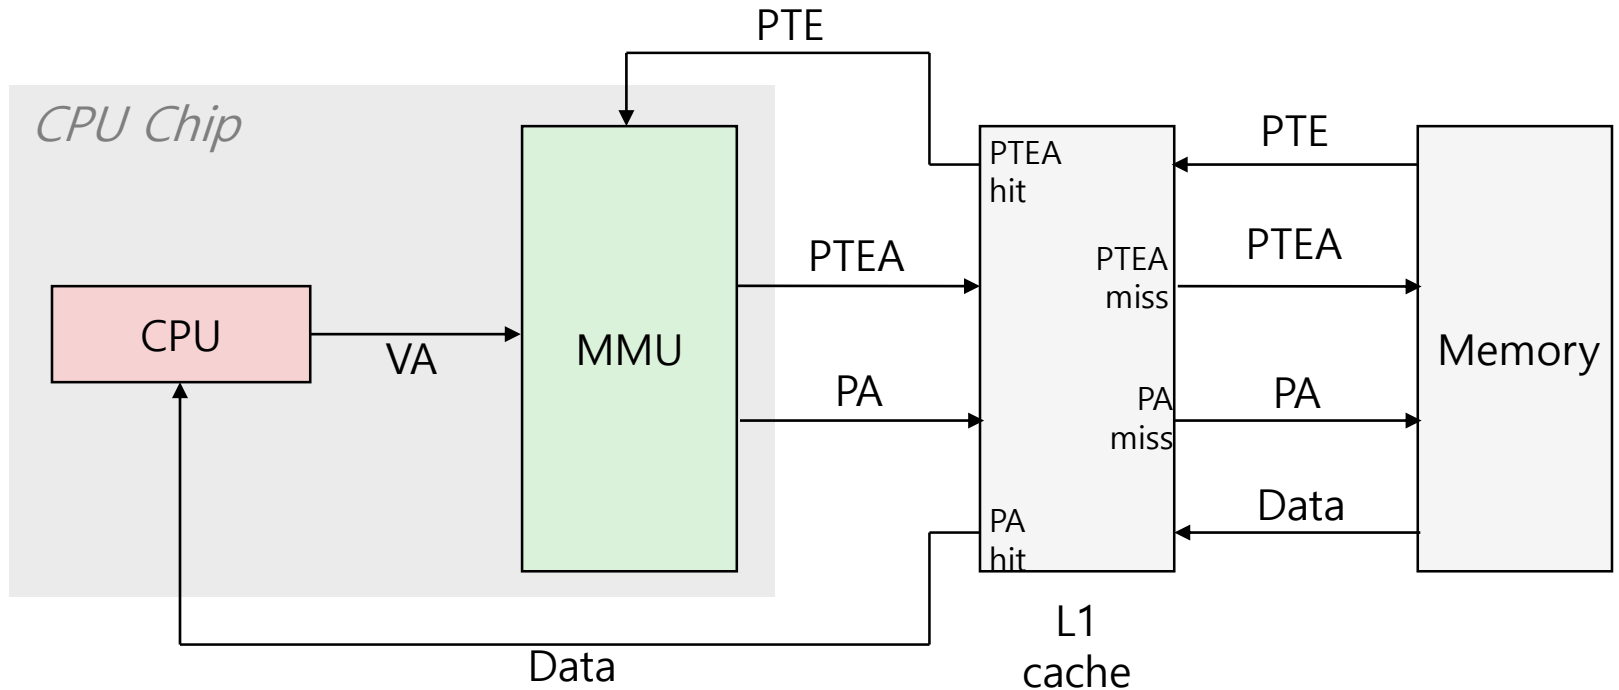
- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor

Address Translation: Page Fault



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Integrating VM and Cache



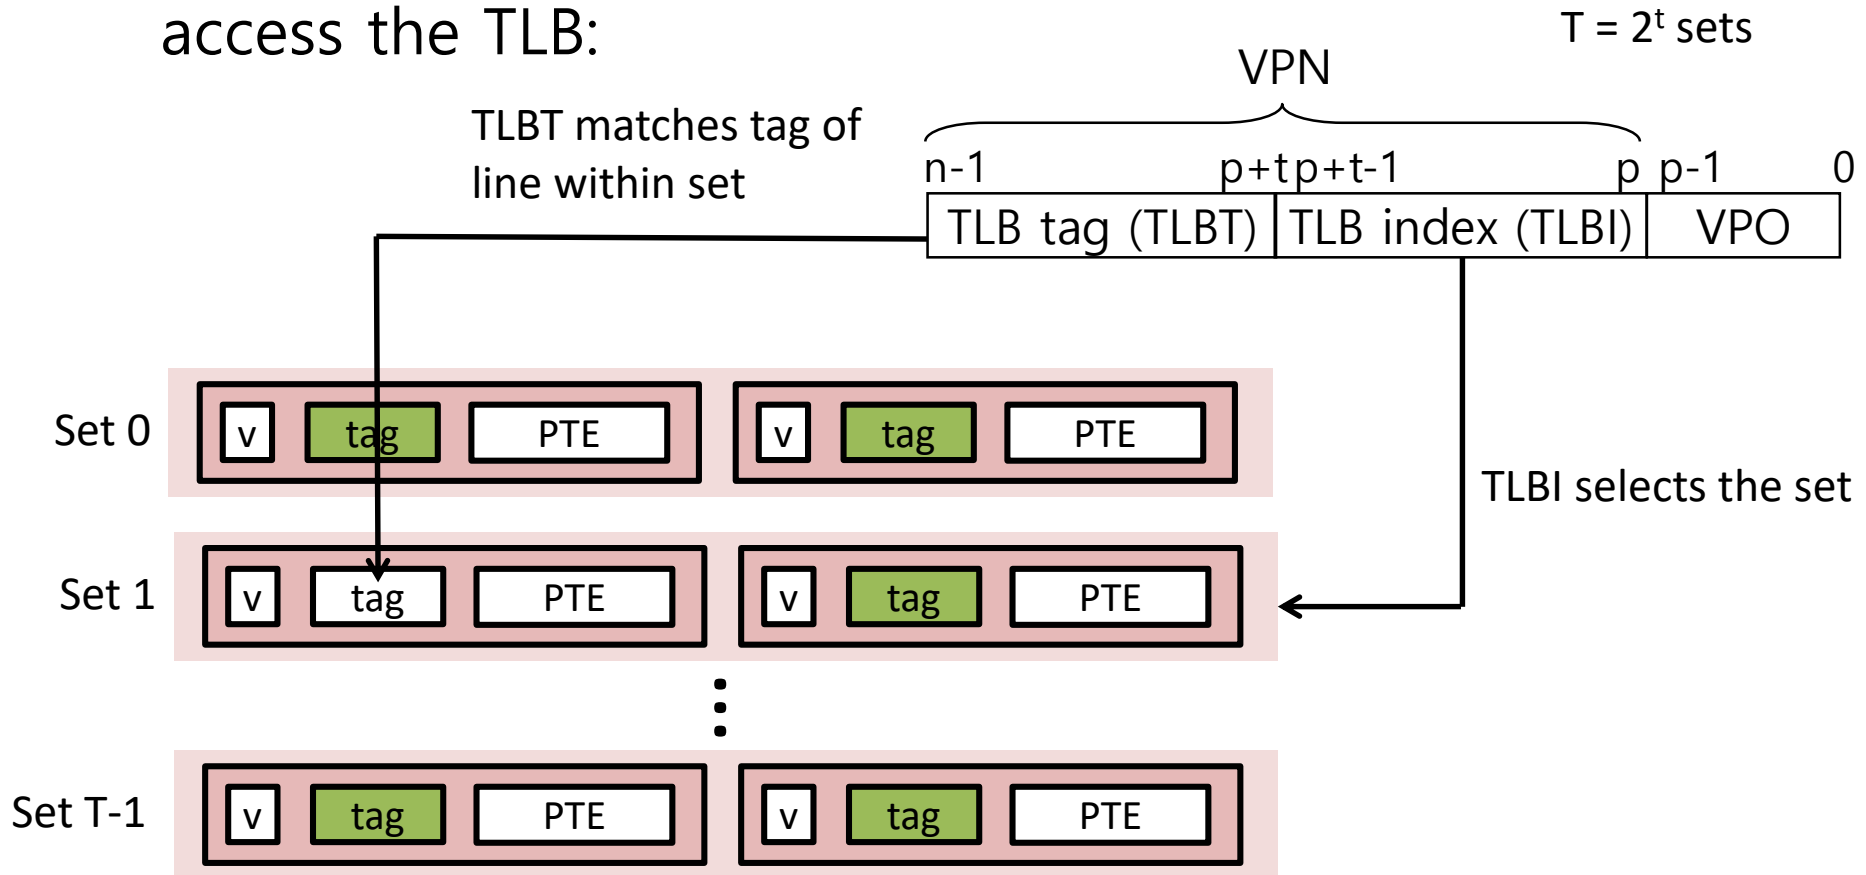
VA: virtual address, PA: physical address, PTE: page table entry, PTEA = PTE address

Speeding up Translation with a TLB

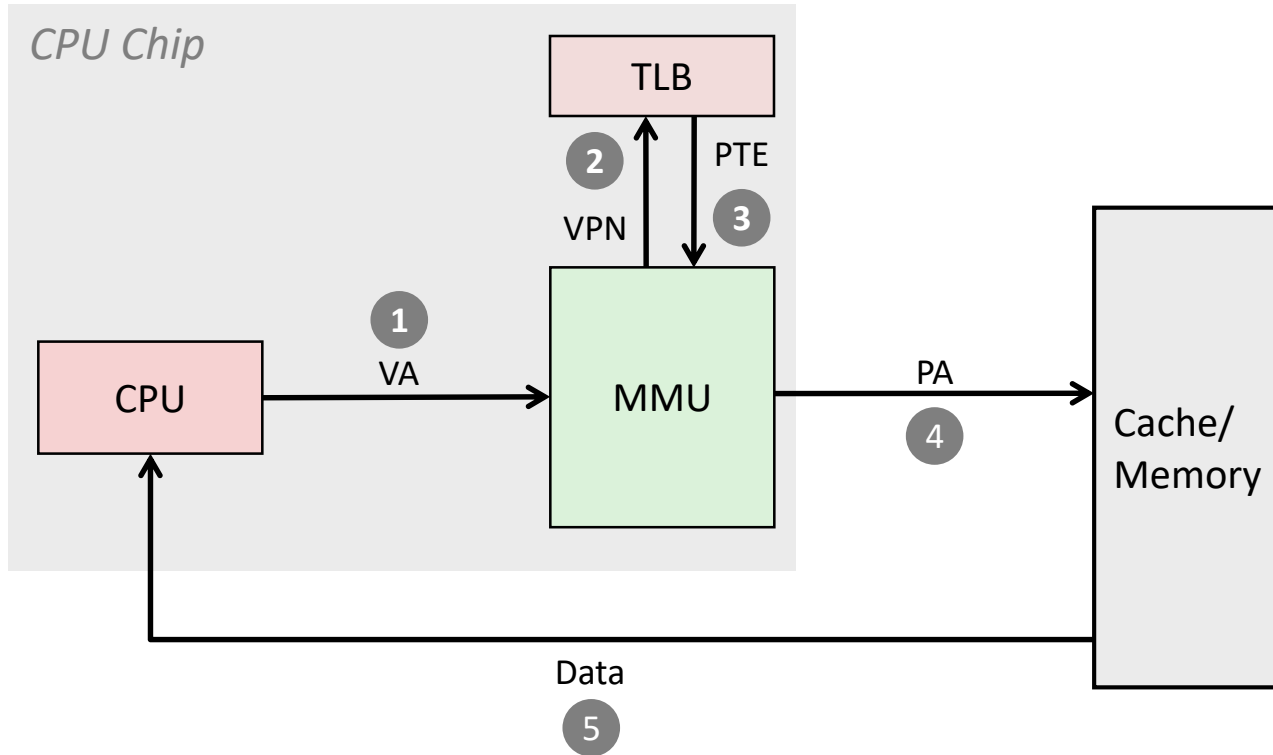
- Page table entries (PTEs) are cached in L1 like any other memory word
 - PTEs may be evicted by other data references
 - PTE hit still requires a small L1 delay
- Solution: *Translation Lookaside Buffer* (TLB)
 - Small hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages

Accessing the TLB

- MMU uses the VPN portion of the virtual address to access the TLB:

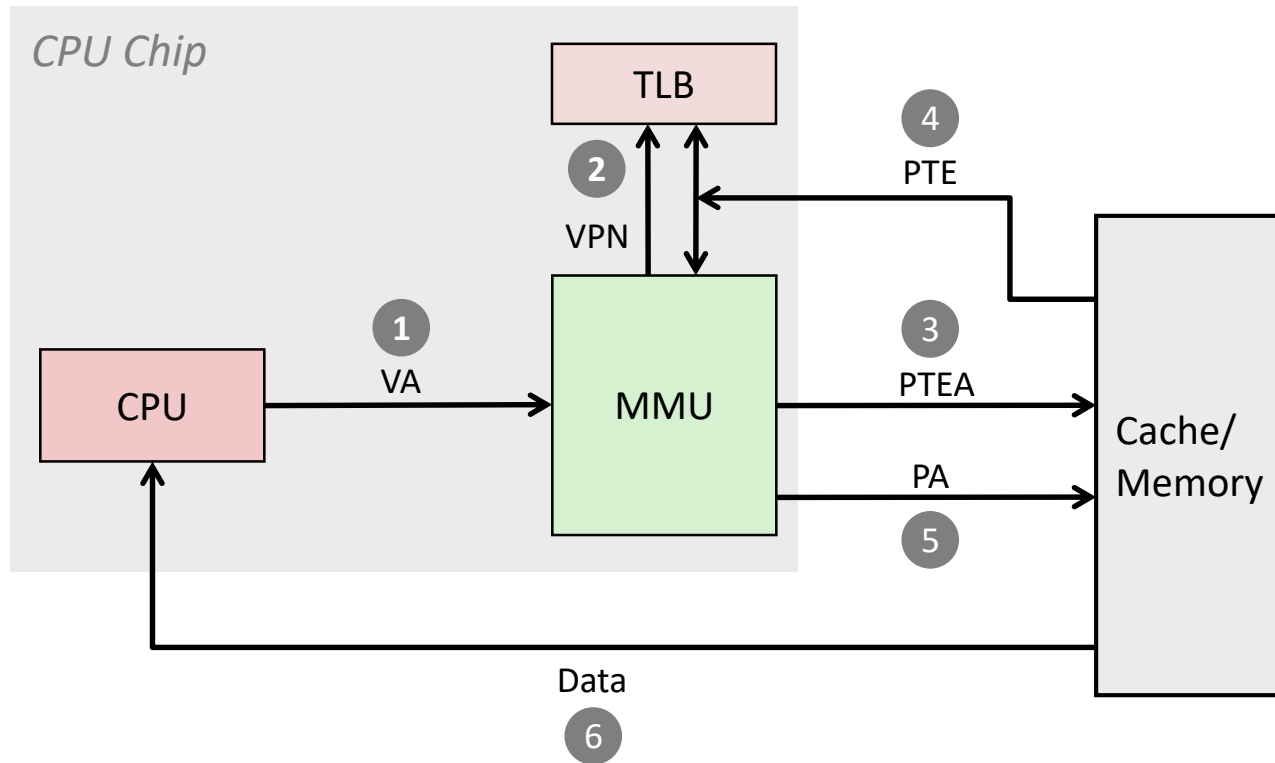


TLB Hit



A TLB hit eliminates a memory access

TLB Miss




A TLB miss incurs an additional memory access (the PTE)

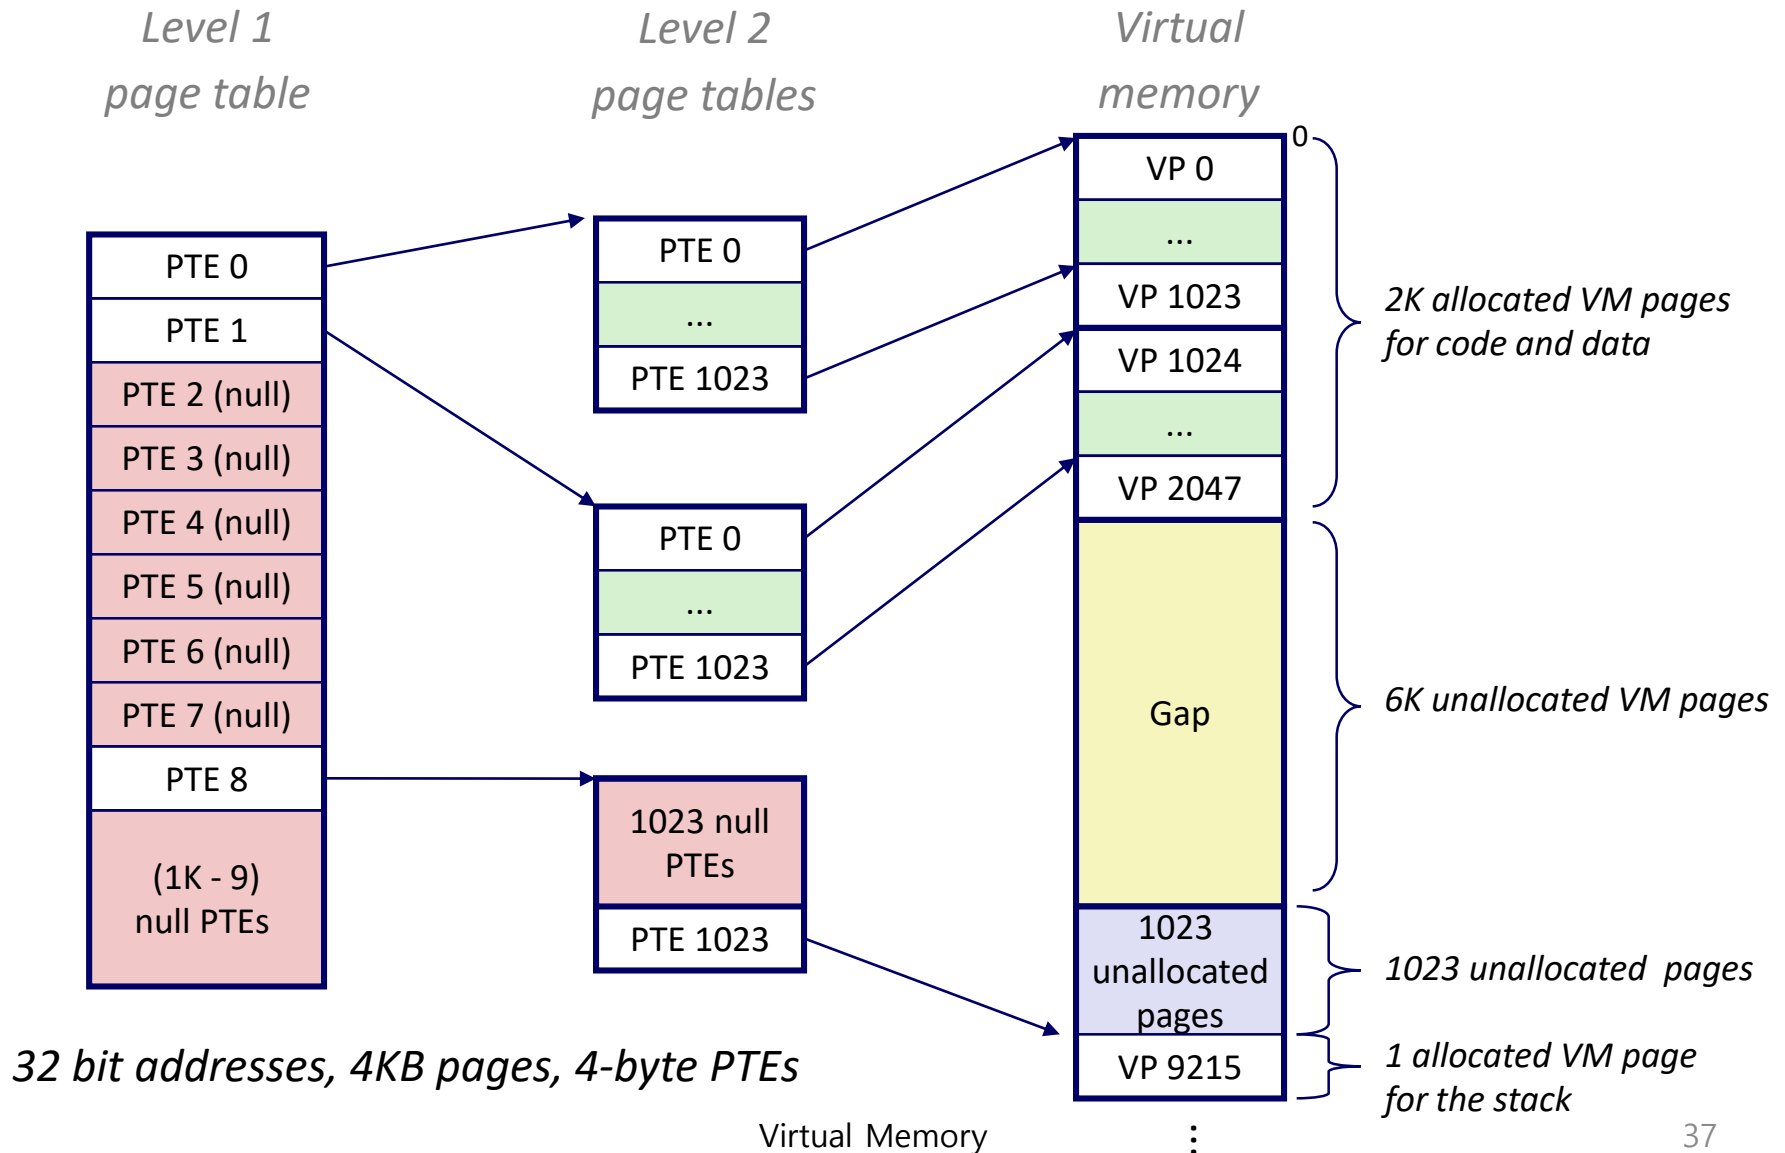
Fortunately, TLB misses are rare. Why?

Multi-Level Page Tables

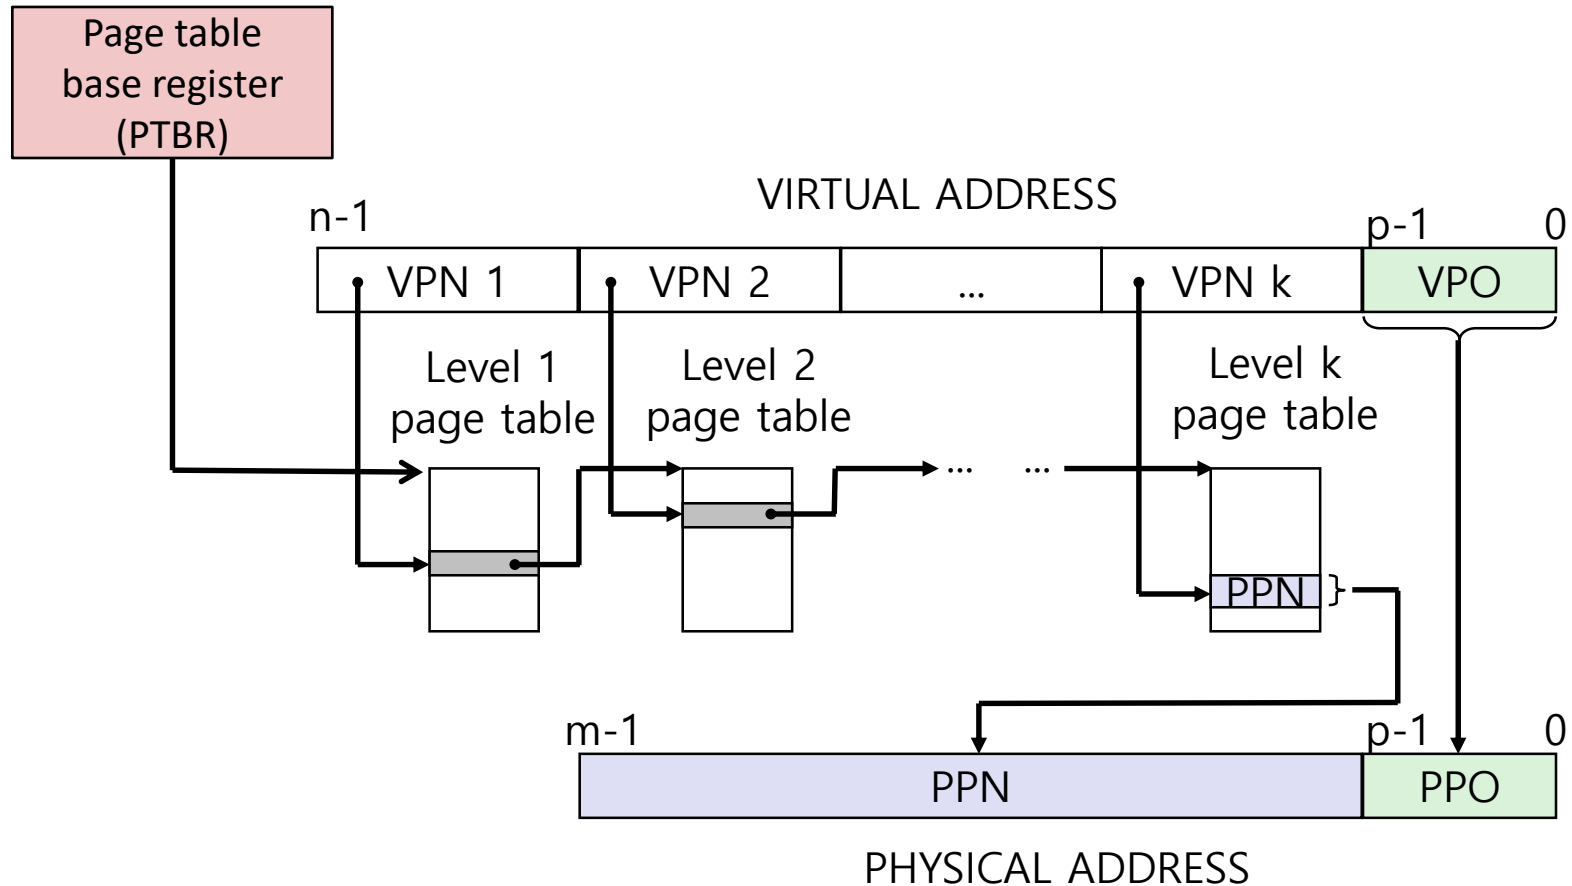
- Suppose:
 - 4KB (2^{12}) page size, 48-bit address space, 8-byte PTE
- Problem:
 - Would need a 512 GB page table!
 - $2^{48} * 2^{-12} * 2^3 = 2^{39}$ bytes
- Common solution: Multi-level page table
- Example: 2-level page table
 - Level 1 table: each PTE points to a page table (always memory resident)
 - Level 2 table: each PTE points to a page (paged in and out like any other data)



A Two-Level Page Table Hierarchy

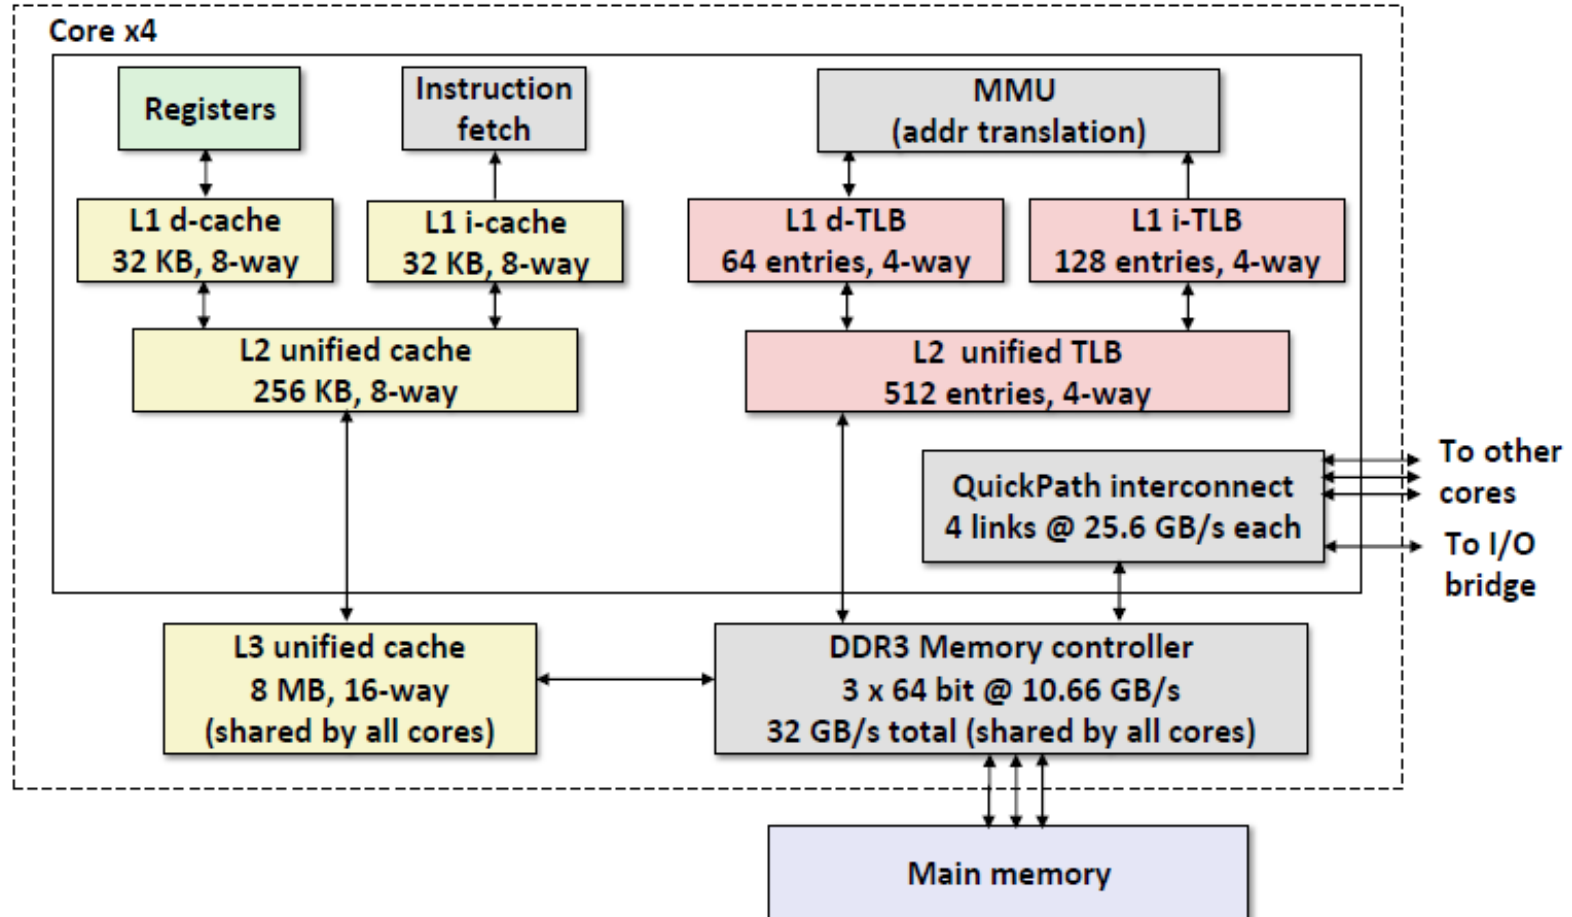


Translating with a k-level Page Table

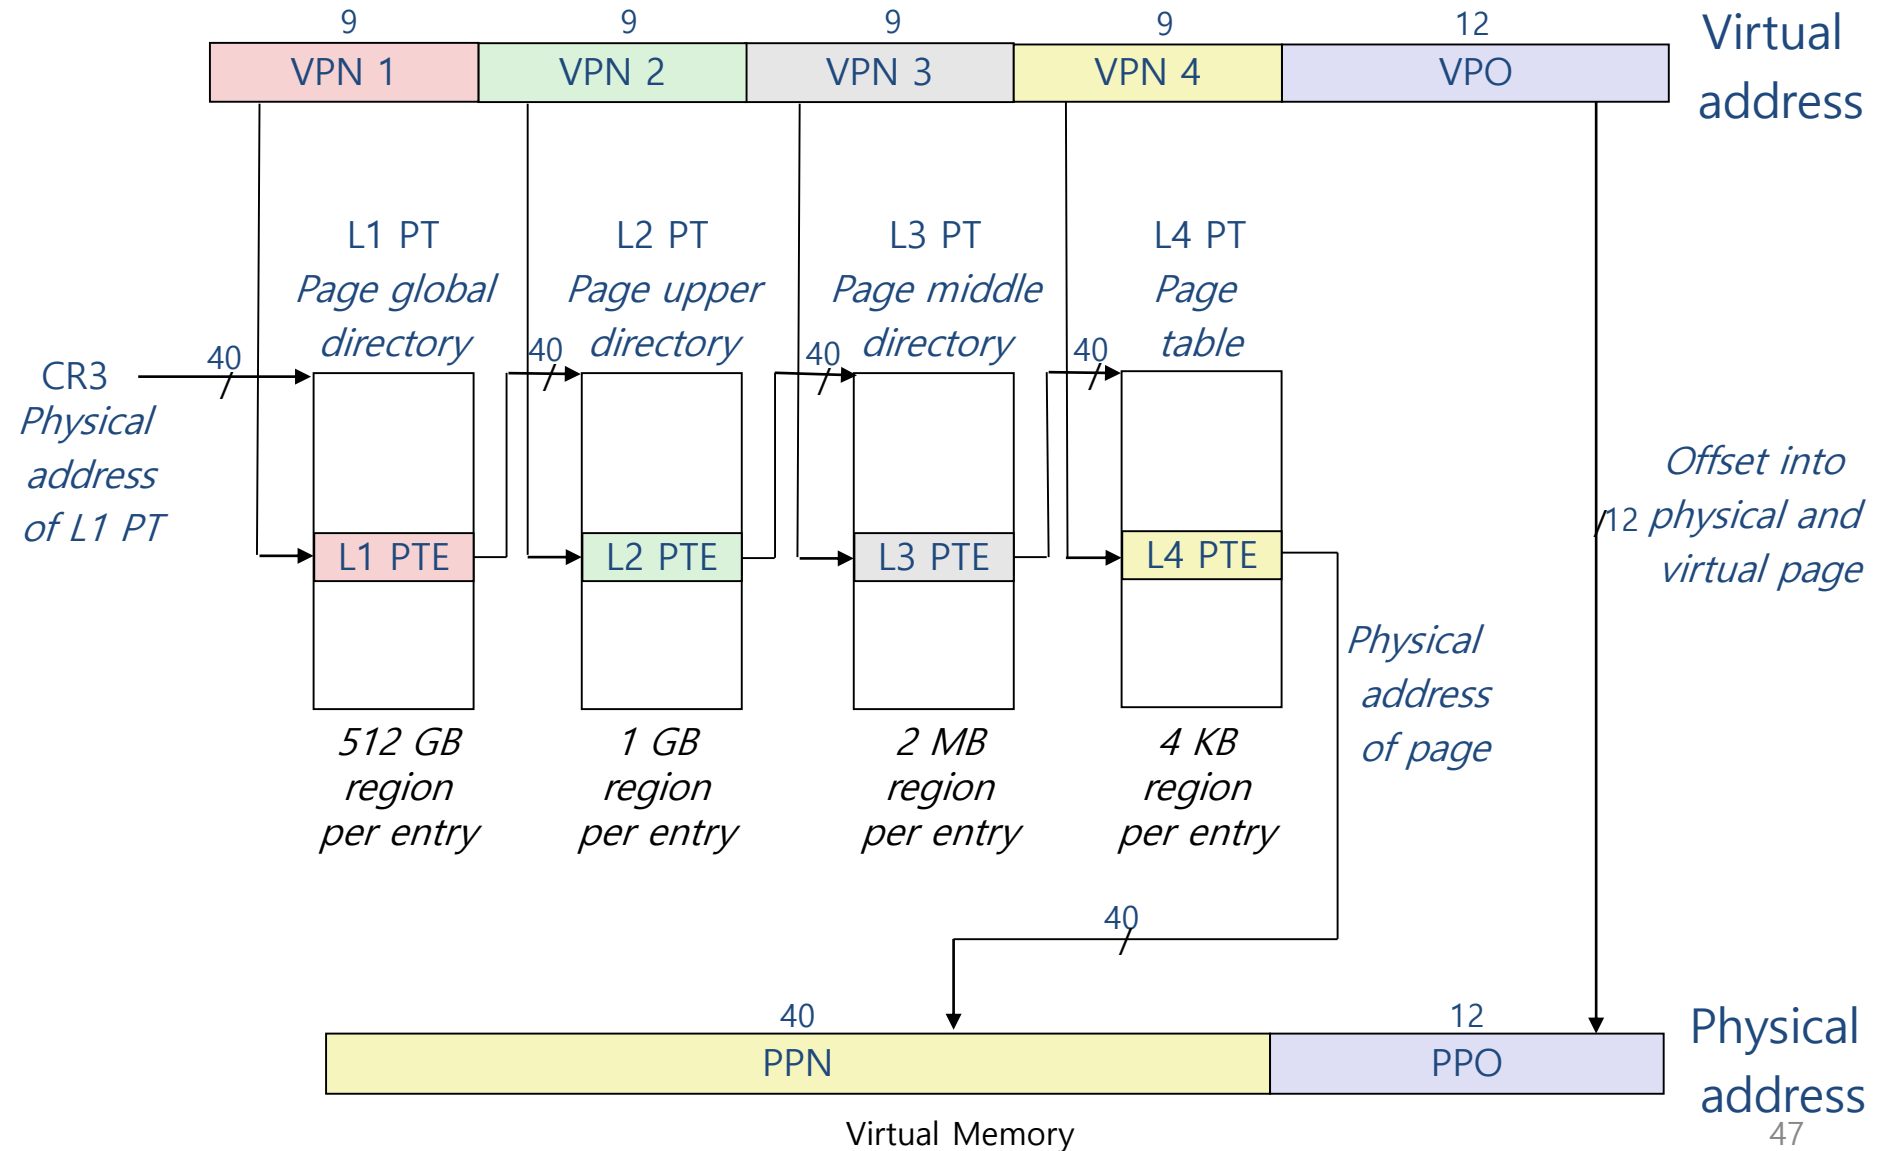


Intel Core i7 memory System

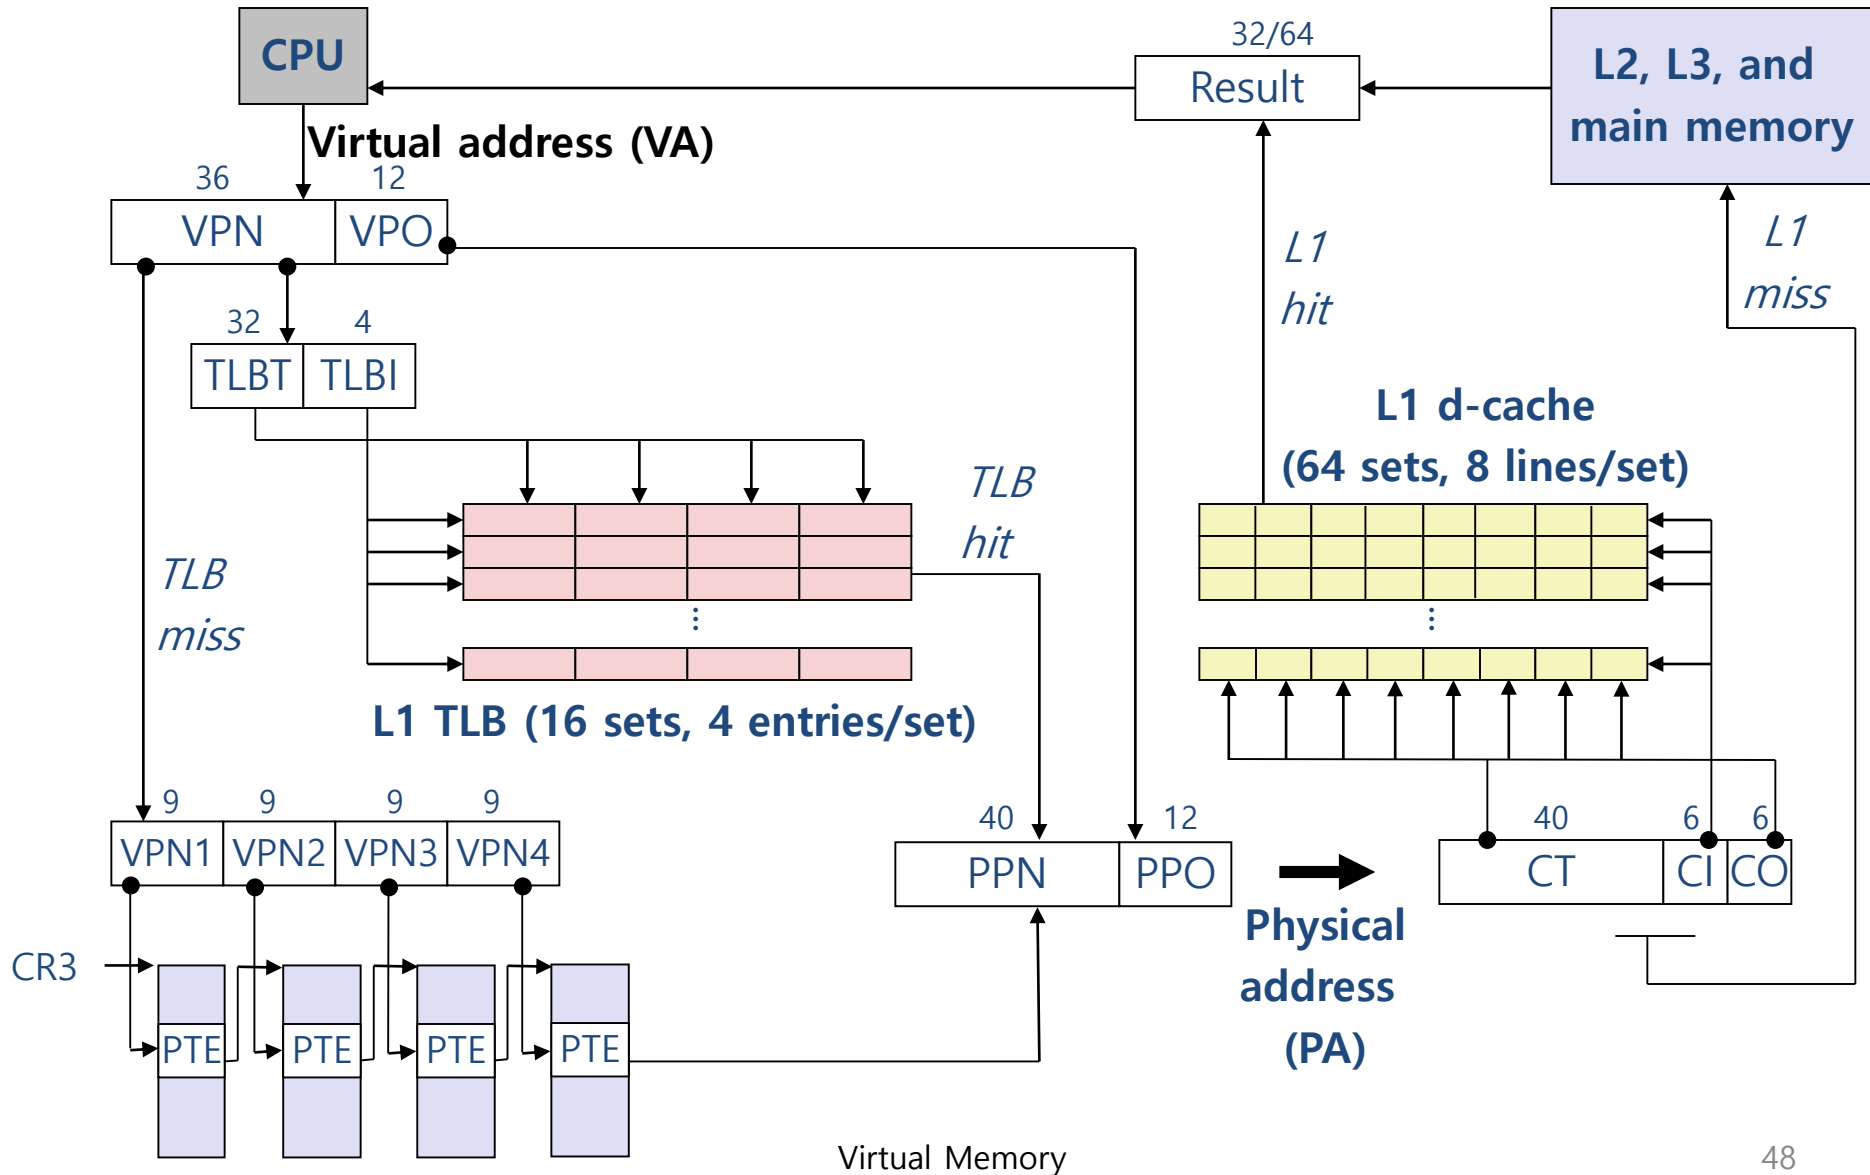
Processor package



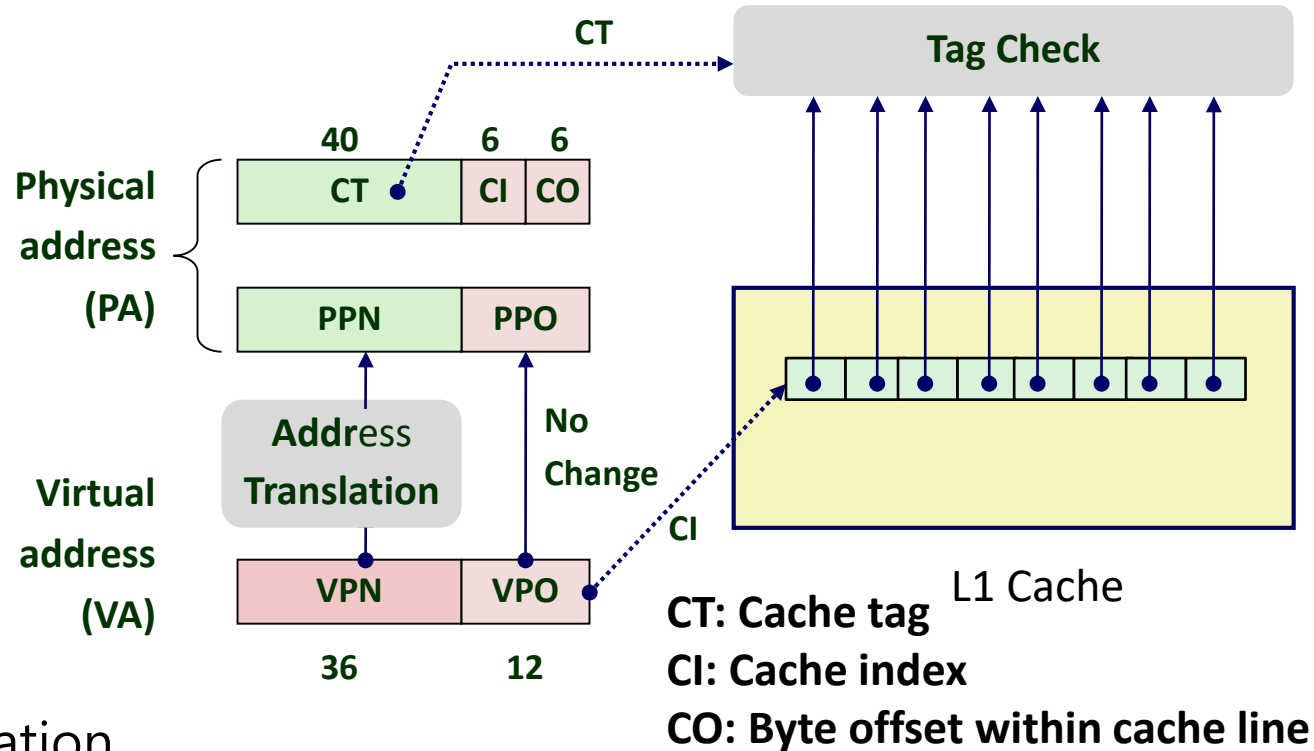
Core i7 Page Table Translation



End-to-end Core i7 Address Translation



Cute Trick for Speeding Up L1 Access



- Observation
 - Bits that determine CI identical in virtual and physical address
 - Can index into cache while address translation taking place
 - Generally we hit in TLB, so PPN bits (CT bits) available next
 - "Virtually indexed, physically tagged"
 - Cache carefully sized to make this possible

Summary

- Programmer's view of virtual memory
 - Each process has its own private linear address space
 - Cannot be corrupted by other processes
- System view of virtual memory
 - Uses memory efficiently by caching virtual memory pages
 - Efficient only because of locality
 - Simplifies memory management and programming
 - Simplifies protection by providing a convenient interpositioning point to check permissions

Questions?